



BITS Pilani
Pilani Campus

COMPUTER ORGANIZATION AND SOFTWARE SYSTEMS

SESSION 11

Dr. Lucy J. Gudino
WILP & Department of CS & IS

Last Session



List of Topic Title	Text/Ref Book/external resource
• Scheduling algorithms (Priority, RR, Multilevel / Feedback)	T2

Today's Session



List of Topic Title	Text/Ref Book/external resource
<ul style="list-style-type: none">• Process Coordination• Synchronization• Deadlock	T2



Process Synchronization


BITS Pilani
Pilani Campus

Introduction




- How do we maximize CPU utilization / improve efficiency?
 - Multiprogramming Vs Multiprocessing
- How to achieve concurrent operation in
 - Uniprocessor System Vs Multiprocessor system

uniprocessor system



multiprocessor system



- Interleaving and overlapping improves processing efficiency, but may produce unpredictable results if not controlled properly
- Example:

```
Procedure echo;  
  Var out,in:Character;  
  Begin  
    input (in, keyboard);  
    out:=in;  
    output(out,Display)  
  End
```

Example (Contd...)



Process Po

1. input (in, keyboard);
2. -----
3. -----
4. -----
5. out:=in;
6. output(out, Display)

Process P1

1. -----
2. input (in, keyboard);
3. out:=in;
4. output(out, Display)

- Reasons unpredictable results :
 - Finite resources
 - Relative speed of execution of processes can not be predicted
 - Sharing of resources (non shareable) among processes
- Processes compete for resources
 - Deadlock
 - Mutual exclusion
 - Starvation
- Cooperating & competing processes can cause problem when executed concurrently → **Synchronization**

Process Synchronization



- Process **Synchronization** means sharing system resources by processes in a such a way that, concurrent access to shared data is handled thereby minimizing the chance of inconsistent data
- A **Critical Section** is a code segment that accesses shared variables or resources and has to be executed as an atomic action.
- Successful use of concurrency requires -
 - Ability to define critical section and
 - Enforce mutual exclusion

Process structure



Process P_i

do {

:

ENTRY SECTION

critical section

Exit SECTION

remainder section

} while (TRUE);

Main requirements



Three requirements

- **A mutual exclusion** : When one process is using a shared modifiable data, the other processes will be excluded from doing the same thing
- **Progress** : when no process in critical section, any process that makes a request is allowed to enter critical section without any delay
- **Bounded Waiting** : Processes requesting critical section should not be delayed indefinitely (no deadlock, starvation)

- ❖ *No assumption should be made about relative execution speed of processes or number of processes*
- ❖ *A process remains inside critical section for a finite amount of time*

Slide 14

Approach to handle Mutual Exclusion



- Software Approach (User is responsible for enforcing Mutual exclusion)
- Hardware Support - Disabling of Interrupt and using Special Instructions
- OS support - Semaphore

Critical Section (Solution1)

innovate

```
Procedure echo;  
  Var out,in:Character;  
Begin  
  input (in, keyboard);  
  out:=in;  
  output(out,Display)  
End
```

Process 0

```
while turn == 1 do {nothing }  
<Critical Section>  
turn = 1
```

Process 1

```
while turn == 0 do {nothing }  
<Critical Section>  
turn = 0
```

"I finished with it, now you have it" → Decker's Algorithm

- Drawback 1: processes must strictly alternate - Pace of execution of one process is determined by pace of execution of other processes
- Drawback 2: if one processes fails other processes is permanently blocked

Main requirements



Three requirements

- **A mutual exclusion** : When one process is using a shared modifiable data, the other processes will be excluded from doing the same thing
- **Progress** : when no process in critical section, any process that makes a request is allowed to enter critical section without any delay
- **Bounded Waiting** : Processes requesting critical section should not be delayed indefinitely (no deadlock, starvation)

- ❖ *No assumption should be made about relative execution speed of processes or number of processes*
- ❖ *A process remains inside critical section for a finite amount of time*

Slide 14

Approach to handle Mutual Exclusion



- Software Approach (User is responsible for enforcing Mutual exclusion)
- Hardware Support - Disabling of Interrupt and using Special Instructions
- OS support - Semaphore

Critical Section (Solution1)

innovate

```
Procedure echo;  
  Var out,in:Character;  
Begin  
  input (in, keyboard);  
  out:=in;  
  output(out,Display)  
End
```

Process 0

```
while turn == 1 do {nothing }  
<Critical Section>  
turn = 1
```

Process 1

```
while turn == 0 do {nothing }  
<Critical Section>  
turn = 0
```

"I finished with it, now you have it" → Decker's Algorithm

- Drawback 1: processes must strictly alternate - Pace of execution of one process is determined by pace of execution of other processes
- Drawback 2: if one processes fails other processes is permanently blocked

Critical Section (Solution 2)

Peterson's solution



- Good algorithmic description of solving the problem
- Two process solution
- The two processes share two variables:
 - `int turn;`
 - `Boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready!

```
Process Pi
do {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j);

    critical section

    flag[i] = false;
    -----
} while (true);
```

Algorithm

turn = 0

flag = {F, F}



t	P0	P1
1	flag[0] = true; turn = 1;	
2		flag[1] = true; turn = 0;
3	while (flag[1] && turn == 1); critical section	
4		while (flag[0] && turn == 0); critical section
5	Critical Section... flag[0] = false	
6		while (flag[0] && turn == 0); critical section flag[1] = false;

Algorithm

turn = 0

flag = {F, F}



t	P0	P1
1		flag[1] = true; turn = 0;
2	flag[0] = true; turn = 1;	
3		while (flag[0] && turn == 0); critical section
4	while (flag[1] && turn == 1); critical section	
5		while (flag[0] && turn == 0); critical section flag[1] = false;
6	Critical Section... flag[0] = false	

Synchronization - Hardware Approach



- Protecting critical regions via locks

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```

Synchronization - Hardware Approach...



- Uniprocessors - could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient
- Modern machines provide special atomic hardware instructions
 - **Atomic** = non-interruptible
 - Example : TestAndSet instruction, Swap instruction

Hardware approach - TestAndSet



- **TestAndSet** instruction used to write to a memory location and return its old value as a single atomic (i.e., non-interruptible) operation.
- **Definition of the TestAndSet () instruction:**

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

Mutual Exclusion with TestAndSet instruction



Shared data: Boolean lock = FALSE;

Process P_0

```
do {  
  while (TestAndSet( &lock));  
          //wait  
  
  critical section  
  lock = FALSE;  
  remainder section  
} while (TRUE);
```

Process P_1

```
do {  
  while (TestAndSet( &lock));  
          //wait  
  
  critical section  
  lock = FALSE;  
  remainder section  
} while (TRUE);
```

```
boolean TestAndSet(boolean *lock)  
{ boolean rv = *lock; *lock = TRUE; return rv; }
```

Slide 24

Semaphores

- is a variable which is treated in a special way
- allow processes to make use of critical section in exclusion to other processes
- process wanting to access critical section locks semaphore and releases lock on exit.
- is a synchronization tool

- Basic properties of semaphore:
 - semaphore S - integer variable with non negative values
 - can only be accessed via two operations
- wait (S):*
- ```
while $S \leq 0$ do no-op;
 $S--$;
```
- signal (S):*
- ```
 $S++$ ;
```
- Semaphore operation is atomic and indivisible
 - wait and signal operations are carried out without interruption

Types of semaphore



- Binary semaphore : can have two values 0 and 1
 - Also known as **mutex locks**, as they are locks that provide *mutual exclusion*.
- *Counting Semaphore* : integer value can range over an unrestricted domain.

Critical Section of n Processes



Shared data:
semaphore S ; *//initially $S = 1$*

wait (S):
 while $S \leq 0$ do *no-op*;
 $S--$;
signal (S):
 $S++$;

Process $P0$:
do {
 remainder section

 wait(S);
 critical
 signal(S);

 remainder section
} while (1);

Process $P1$:
do {
 remainder section

 wait(S);
 critical section
 signal(S);

 remainder section
} while (1);

Slide 28

Process synchronization



- Semaphore can be used to solve various synchronization problems
- Example : two concurrent processes : P1 (with statement S1) and P2 (with statement S2)
 - requirement : **s2 should be executed only after s1 is executed**
 - semaphore variable : **synch** initialized to zero.

P1:

S1;

signal (synch);

P2:

wait (synch);

S2;

wait (synch):

while $\text{synch} \leq 0$ do *no-op*;

$\text{synch}--$;

signal (synch):

$\text{synch}++$;

Slide 29

- Main disadvantage : Busy Waiting → waiting wastes CPU cycles
- semaphore is also called a **spinlock** because the process "spins" while waiting for the lock.
- Solution: on finding zero semaphore value (binary semaphore), the process can block itself instead of busy waiting
- The block operation places a process into a waiting queue associated with the semaphore, and the state of the process is switched to the waiting state.
- Then control is transferred to the CPU scheduler, which selects another process to execute.

- A process that is blocked, waiting on a semaphore S , should be restarted when some other process executes a `signal()` operation.
- The process is restarted by a `wakeup ()` operation, which changes the process from the waiting state to the ready state.
- The process is then placed in the ready queue.

Consumer – Producer Problem

- Also known as bounded buffer problem
- Two processes – consumer and producer
- Consumer and Producer processes share a common, fixed size buffer
- Producer process : generates data and puts it in the buffer
- Consumer process: consumes data from the buffer
- *Problem statement : When a producer is placing an item in the buffer, then at the same time consumer should not consume any item.*

$\text{mutex} = 1$

$\text{Full} = 0 \rightarrow$ Initially, all slots are empty.

$\text{Empty} = n \rightarrow$ All slots are empty



Producer

```
do{
//produce an item
wait(empty);
wait(mutex);
    //place in buffer
signal(mutex);
signal(full);

}while(true)
```

Consumer

```
do{
wait(full);
wait(mutex);
    // remove item from buffer
signal(mutex);
signal(empty);
    // consumes item

}while(true)
```

Deadlock

innovate

achieve

lead

EXAMPLES:

"It takes money to make money".

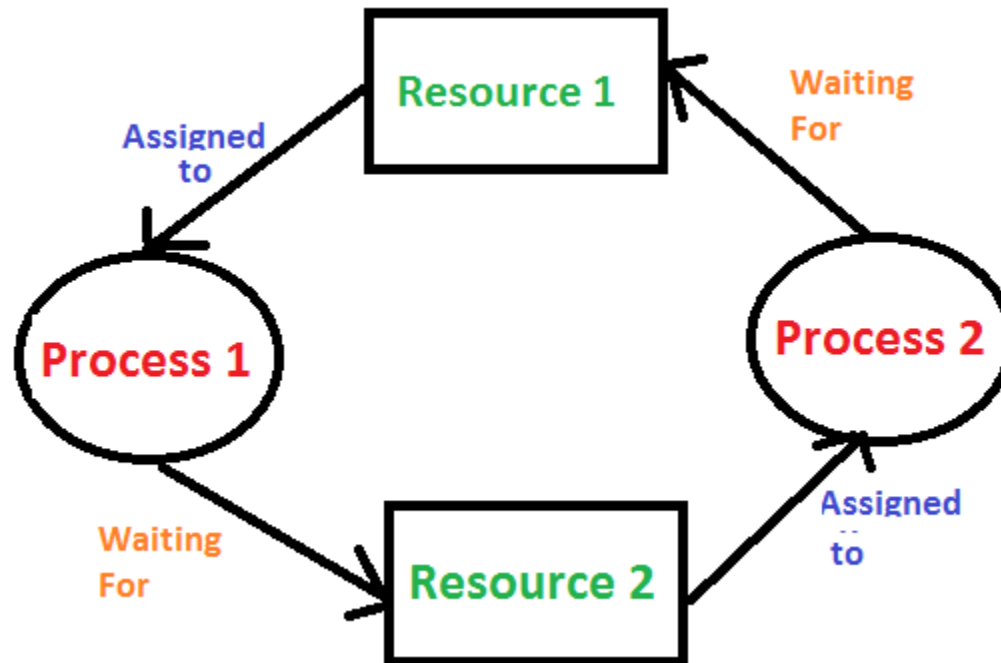
"You can't get a job without experience; you can't get experience without a job."



Slide 34

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.



- Example
 - System has 2 disk drives D1 and D2
 - P_1 and P_2 each hold one disk drive and each needs another one.
- Finite number of resources
 - Example : Memory space, CPU, files, I/O devices - printers, monitor, DVD drives
- Resource types and instances
 - system with 3 printers → Resource Type: printer and Number of instances : 3
- Process must request a resource before using it and must release the resource after using it

- Each process utilizes a resource as follows:
 - request
 - use
 - release
- Device : request () and release()
- File: open() and close ()
- Memory: allocate () and free ()

Deadlock Characterization


Mutual exclusion: only one process at a time can use a resource.

Hold and wait: a process holding at least one resource and is waiting to acquire additional resources held by other processes.

No preemption: a resource can be released only voluntarily by the process holding it, after that process has completed its task.

Circular wait: there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by

P_2, \dots, P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .



NECESSARY CONDITIONS

ALL of these four **must** happen simultaneously for a deadlock to occur Slide 38

Resource-Allocation Graph



A set of vertices V and a set of edges E .

V is partitioned into two types:


- $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
- $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.

request edge - directed edge $P_i \rightarrow R_j$

assignment edge - directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


Process




Resource Type with 4 instances




P_i requests instance of R_j



P_i is holding an instance of R_j



Example 1



R3 Assigned to P3

P2 Requests R3

Slide 41

Example 2

Cycle 1:

$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

Cycle 2:

$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$

