

innovate

achieve

lead



BITS Pilani
Pilani Campus

Machine Learning

DSECL ZG565

Dr. Monali Mavani

Session Content



- Decision Theory - Bishop
- Probabilistic Generative Model versus Probabilistic Discriminative Model
- Logistic Regression –Bayesian Analysis (New chapter Tom Mitchell)
- Estimating Parameters Linear Regression : Closed form solution
- Linear basis function models (3.1 Bishop)
- Evaluation metrics

Decision Theory

Decision Theory

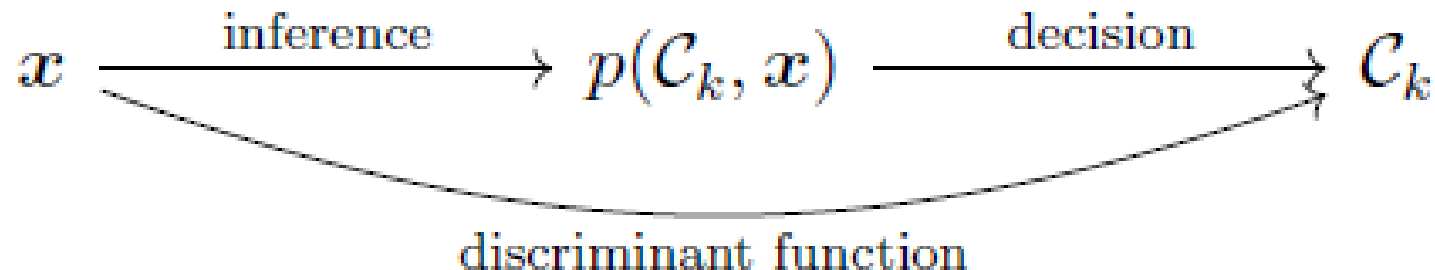


We have broken the **classification problem** down into two separate stages

- **Inference step:** Determine $p(\mathbf{x}, t)$ OR $p(\mathbf{x}/C_k)$ from training data. (**For regression problems t is continuous variables**, whereas for classification problems t represents class labels.)



$$t = \arg \max_k \{ \underbrace{p(\mathbf{x}|C_k)}_{\text{Likelihood}} \}$$

- **Decision step:** Determine optimal t for test input x : how to make optimal decisions given the appropriate probabilities.




Minimum Misclassification Rate

- Divide the input space into regions R_k called decision regions, one for each class, such that all points in R_k are assigned to class C_k
- Boundaries between decision regions are called decision boundaries or decision surfaces
- A mistake occurs when an input vector belonging to class C_1 is assigned to class C_2 or vice versa.



Minimum Misclassification Rate



sum of red and green is the first integral of equation

\hat{x} : decision boundary.
 x_0 : optimal decision boundary

$$x_0 : \arg \min_{\mathcal{R}_1} \{p(\text{mistake})\}$$

- Minimum error decision rule
- To minimize $p(\text{mistake})$ we should arrange that each \mathbf{x} is assigned to whichever class has the smaller value of the integrand
- if $p(\mathbf{x}, C_1) > p(\mathbf{x}, C_2)$ for a given value of \mathbf{x} , then we should assign that \mathbf{x} to class C_1 .
- Since $p(\mathbf{x}, C_k) = p(C_k/\mathbf{x})p(\mathbf{x})$, choose class for which *a posteriori* probability is highest
 - Called Bayes Classifier

$$\begin{aligned} p(\text{mistake}) &= p(\mathbf{x} \in \mathcal{R}_1, \mathcal{C}_2) + p(\mathbf{x} \in \mathcal{R}_2, \mathcal{C}_1) \\ &= \int_{\mathcal{R}_1} p(\mathbf{x}, \mathcal{C}_2) \, d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, \mathcal{C}_1) \, d\mathbf{x}. \end{aligned}$$

Minimum Misclassification Rate

General case of K classes

$$\begin{aligned} p(\text{correct}) &= \sum_{k=1}^K p(\mathbf{x} \in \mathcal{R}_k, \mathcal{C}_k) \\ &= \sum_{k=1}^K \int_{\mathcal{R}_k} p(\mathbf{x}, \mathcal{C}_k) d\mathbf{x} \end{aligned}$$



Three distinct approaches to Decision Problems

1. Generative
2. Discriminative
3. Discriminant Function

1. Generative Models



- Model class-conditional pdfs and prior probabilities
- First solve inference problem of determining class-conditional densities $P(X|Y)$, for each class separately
- Then use Bayes theorem to determine posterior probabilities
- Then use decision theory to determine class membership
- “Generative” since sampling can generate synthetic data points
- Use the capacity of the model to characterize how the data is generated (both inputs and outputs)
- explicitly models the **actual distribution of each class**.
- Gaussians, Naïve Bayes, Mixtures of multinomials , Mixtures of Gaussians, Bayesian networks,

2. Discriminative Models

- Directly estimate posterior probabilities $P(Y|X)$ directly from training data and use decision theory to assign each new x to one of the classes
- No attempt to model underlying probability distributions
- models the **decision boundary between the classes**
- Logistic regression, SVMs , **tree based classifiers (e.g. decision tree)** Traditional neural networks, Nearest neighbor

3. Discriminant Functions

- Find a function $f(x)$ that maps each input x
- directly to class label
 - In two-class problem, $f(.)$ is binary valued
 - $f=0$ represents class $C1$ and $f=1$ represents class $C2$
- Probabilities play no role
 - No access to posterior probabilities $p(Y|X)$
 - Linear discriminant, Fisher Linear Disc, Perceptron

Naive Bayes and logistic regression: two different modelling paradigms



Spam classification problem

First Strategy: discriminative (e.g., logistic regression)

- Use training set to find a decision boundary in the feature space that separates spam and non-spam emails
- Given a test point, predict its label based on which side of the boundary it is on.


Second Strategy: generative (e.g., naive bayes)

- Look at spam emails and build a model of what they look like.
- Similarly, build a model of what non-spam emails look like.
- To classify a new email, match it against both the spam and non-spam models to see which is the better fit.


Probabilistic Generative Model versus Probabilistic Discriminative Model



Generative:



Discriminative:



Probabilistic Generative Model versus Probabilistic Discriminative Model



Generative	Discriminative
Ex: Naïve Bayes	Ex: Logistic Regression
Estimate $P(Y)$ and $P(X Y)$	Finds class label directly $P(Y X)$
Prediction $\hat{y} = \operatorname{argmax}_y P(Y = y)P(X = x Y = y)$	Prediction $\hat{y} = P(Y = y X = x) = \frac{1}{1+e^{-\theta^\top x}}$
Decision boundary	Probability distributions of the data

Logistic Regression –Bayesian Analysis

Logistic Regression and Gaussian Naïve Bayes Classifier



- Interestingly, the parametric form of $P(Y|X)$ used by Logistic Regression is precisely the form implied by the assumptions of a Gaussian Naive Bayes classifier.
- Therefore, we can view Logistic Regression as a closely related alternative to GNB, though the two can produce different results in many cases



- Estimation of $n+1$ parameters say, $\theta = (\theta_0, \theta_1, \dots, \theta_n)$ for which the probability of the observed data is the maximum on the training data set.
 - generative learning
 - the Gaussian Naive Bayes assumptions
 - discriminative learning
 - maximize the conditional data likelihood (or minimizing negative likelihood)
 - Uses gradient descent optimization



Bayesian inference for logistic analyses follows the usual pattern for all Bayesian analyses:

1. Write down the likelihood function of the data
2. Form a prior distribution over all unknown parameters.
3. Use Bayes theorem to find the posterior distribution over all parameters.

Where does the **hypothesis function** come from?

- Logistic regression hypothesis representation

$$P(Y=1|X) = h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}} = \frac{1}{1+e^{-(\theta_0+\theta_1 x_1+\theta_2 x_2+\dots+\theta_n x_n)}}$$

- Consider learning $f: X \rightarrow Y$, where
 - X is a vector of real-valued features $[X_1, \dots, X_n]^T$
 - Y is Boolean
 - Assume all X_i are conditionally independent given Y
 - Model likelihood $P(X_i|Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i)$ and assume variance is independent of class, i.e. $\sigma_{i0} = \sigma_{i1} = \sigma_i$**

$$P(X_i|Y = y_k) = \frac{1}{\sqrt{2\pi\sigma_{ik}^2}} e^{-\frac{1}{2}\left(\frac{X_i - \mu_{ik}}{\sigma_{ik}}\right)^2}$$

$$P(x|y_k) = \frac{1}{\sqrt{2\pi}\sigma_i} e^{-\frac{(x-\mu_{ik})^2}{2\sigma_i^2}}$$

- Model prior $P(Y)$ as Bernoulli π : $P(Y=1) = \pi$ and $P(Y=0) = 1-\pi$**

What is $P(Y|X_1, X_2, \dots, X_n)$?

Logistic Regression –Bayesian Analysis

innovate

achieve

lead

$$P(Y = 1|X) = \frac{P(Y = 1)P(X|Y = 1)}{P(Y = 1)P(X|Y = 1) + P(Y = 0)P(X|Y = 0)}$$

Applying Bayes rule

$$P(Y = 1|X) = \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}}$$

Divide by $P(Y = 1)P(X|Y = 1)$

$$P(Y = 1|X) = \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})}$$

Apply $\exp(\ln(\cdot))$

By independence assumption:

$$\frac{P(X|Y = 0)}{P(X|Y = 1)} = \prod_i \frac{P(X_i|Y = 0)}{P(X_i|Y = 1)}$$

$P(Y=1)=\pi$ and $P(Y=0)=1-\pi$
by modelling $P(Y)$ as Bernoulli

$$P(Y = 1|X) = \frac{1}{1 + \exp(\ln \frac{P(Y = 0)}{P(Y = 1)} + \ln \frac{P(X|Y = 0)}{P(X|Y = 1)})}$$

$$P(Y = 1|X) = \frac{1}{1 + \exp(\ln \frac{1-\pi}{\pi} + \sum_i \ln \frac{P(X_i|Y = 0)}{P(X_i|Y = 1)})}$$

$$P(Y = 1|X) = \frac{1}{1 + \exp(\ln \frac{1-\pi}{\pi} + \ln \prod_i \frac{P(X_i|Y = 0)}{P(X_i|Y = 1)})}$$

Logistic Regression –Bayesian Analysis



Plug in $P(X_i|Y)$

$$P(Y = 1|X) = \frac{1}{1 + \exp(\ln \frac{1-\pi}{\pi} + \sum_i \left(\frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2} \right))}$$

$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)} = \frac{1}{1 + \exp(\theta_0 + \sum_i \theta_i X_i)}$$

$$w_0 = \ln \frac{1-\pi}{\pi} + \sum_i \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2}$$

$$w_i = \frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2}$$

$$P(Y = 0|X) = 1 - P(Y = 1|X) = \frac{\exp(w_0 + \sum_{i=1}^n w_i X_i)}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

linear classification rule

$$\frac{P(Y = 0|X)}{P(Y = 1|X)} = \exp(w_0 + \sum_i w_i X_i)$$

$$\ln \frac{P(Y = 0|X)}{P(Y = 1|X)} = w_0 + \sum_i w_i X_i$$

$P(X_i|Y)$ derivation

innovate

achieve

lead

$$\begin{aligned}\sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)} &= \sum_i \ln \frac{\frac{1}{\sqrt{2\pi\sigma_i^2}} \exp\left(\frac{-(X_i-\mu_{i0})^2}{2\sigma_i^2}\right)}{\frac{1}{\sqrt{2\pi\sigma_i^2}} \exp\left(\frac{-(X_i-\mu_{i1})^2}{2\sigma_i^2}\right)} \\&= \sum_i \ln \exp\left(\frac{(X_i-\mu_{i1})^2 - (X_i-\mu_{i0})^2}{2\sigma_i^2}\right) \\&= \sum_i \left(\frac{(X_i-\mu_{i1})^2 - (X_i-\mu_{i0})^2}{2\sigma_i^2}\right) \\&= \sum_i \left(\frac{(X_i^2 - 2X_i\mu_{i1} + \mu_{i1}^2) - (X_i^2 - 2X_i\mu_{i0} + \mu_{i0}^2)}{2\sigma_i^2}\right) \\&= \sum_i \left(\frac{2X_i(\mu_{i0} - \mu_{i1}) + \mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2}\right) \\&= \sum_i \left(\frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2}\right)\end{aligned}$$

$$P(Y=1|X) = \frac{1}{1 + \exp(\ln \frac{1-\pi}{\pi} + \sum_i \left(\frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2}\right))}$$



Where does the **cost** come from? - Logistic regression

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

Learning: fit parameter θ

$$\min_{\theta} J(\theta)$$

Prediction: given new x

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Slide credit: Andrew Ng

maximum conditional likelihood - θ_{MCLE}



- **Goal:** choose θ to maximize conditional likelihood of training data

$$- P_{\theta}(Y = 1|X = x) = h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$

$$- P_{\theta}(Y = 0|X = x) = 1 - h_{\theta}(x) = \frac{e^{-\theta^T x}}{1+e^{-\theta^T x}}$$

conditional likelihood

- **Training data** $D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
- **Data likelihood** $= \prod_{i=1}^m P_{\theta}((x^{(i)}, y^{(i)}))$
- **Data conditional likelihood** $= \prod_{i=1}^m P_{\theta}(y^{(i)}|x^{(i)})$

$$\theta_{\text{MCLE}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^m P_{\theta}(y^{(i)}|x^{(i)})$$

Slide credit: Tom Mitchell

Expressing conditional log-likelihood

$$L(\theta) = \log \prod_{i=1}^m P_{\theta}(y^{(i)} | x^{(i)}) = \sum_{i=1}^m \log P_{\theta}(y^{(i)} | x^{(i)})$$

Recall, each label $y^{(i)}$ is binary with prob. $P_{\theta}(y^{(i)} | x^{(i)})$: Assume Bernoulli likelihood: (use PMF of Bernoulli dist.)

$$\begin{aligned} &= \sum_{i=1}^m y^{(i)} \log P_{\theta}(y^{(i)} = 1 | x^{(i)}) (1 - y^{(i)}) \log P_{\theta}(y^{(i)} = 0 | x^{(i)}) \\ &= \sum_{i=1}^m y^{(i)} \log (h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \end{aligned}$$

To turn Loglikelihood into loss function flip the sign

$$\theta_{MCLE} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^m P_{\theta}(y^{(i)} | x^{(i)}) = - \underset{\theta}{\operatorname{argmin}} \prod_{i=1}^m P_{\theta}(y^{(i)} | x^{(i)})$$

$$\theta_{MCLE} = - \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^m y^{(i)} \log (h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

Learning Model Parameters – Closed Form Solution (using vectorization)

Vectorization

- Benefits of vectorization
 - More compact equations
 - Faster code (using optimized matrix libraries)

- Consider our model:

$$h(\mathbf{x}) = \sum_{j=0}^d \theta_j x_j$$

- Let

$$\boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix} \quad \mathbf{x}^\top = \begin{bmatrix} 1 & x_1 & \dots & x_d \end{bmatrix}$$

- Can write the model in vectorized form as $h(\mathbf{x}) = \boldsymbol{\theta}^\top \mathbf{x}$

Vectorization

- Consider our model for n instances:

$$h\left(\mathbf{x}^{(i)}\right)=\sum_{j=0}^d \theta_j x_j^{(i)}$$

- Let

$$\boldsymbol{\theta}=\left[\begin{array}{c} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{array}\right] \quad \mathbf{X}=\left[\begin{array}{cccc} 1 & x_1^{(1)} & \cdots & x_d^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & \cdots & x_d^{(i)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(n)} & \cdots & x_d^{(n)} \end{array}\right]$$

$\mathbb{R}^{(d+1) \times 1}$
 $\mathbb{R}^{n \times (d+1)}$

- Can write the model in vectorized form as $h_{\boldsymbol{\theta}}(\mathbf{x})=\mathbf{X} \boldsymbol{\theta}$

Vectorization

- For the linear regression cost function:

$$J(\boldsymbol{\theta}) = \frac{1}{2n} \sum_{i=1}^n \left(h_{\boldsymbol{\theta}} \left(\mathbf{x}^{(i)} \right) - y^{(i)} \right)^2$$

$$= \frac{1}{2n} \sum_{i=1}^n \left(\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

$$= \frac{1}{2n} \underbrace{(\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T}_{\mathbb{R}^{1 \times n}} \underbrace{(\mathbf{X}\boldsymbol{\theta} - \mathbf{y})}_{\mathbb{R}^{n \times 1}}$$

$\mathbb{R}^{n \times (d+1)}$
 $\mathbb{R}^{(d+1) \times 1}$


Let:

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

Closed Form Solution

- Instead of using GD, solve for optimal θ analytically

– Notice that the solution is when $\frac{\partial}{\partial \theta} J(\theta) = 0$



- Derivation:

$$\begin{aligned}
 \mathcal{J}(\theta) &= \frac{1}{2n} (\mathbf{X}\theta - \mathbf{y})^\top (\mathbf{X}\theta - \mathbf{y}) \\
 &\propto \theta^\top \mathbf{X}^\top \mathbf{X} \theta - \boxed{\mathbf{y}^\top \mathbf{X} \theta} - \boxed{\theta^\top \mathbf{X}^\top \mathbf{y}} + \mathbf{y}^\top \mathbf{y} \\
 &\propto \theta^\top \mathbf{X}^\top \mathbf{X} \theta - 2\theta^\top \mathbf{X}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y}
 \end{aligned}$$

1 x 1

Take derivative and set equal to 0, then solve for θ :

$$\frac{\partial}{\partial \theta} (\theta^\top \mathbf{X}^\top \mathbf{X} \theta - 2\theta^\top \mathbf{X}^\top \mathbf{y} + \cancel{\mathbf{y}^\top \mathbf{y}}) = 0$$

$$(\mathbf{X}^\top \mathbf{X})\theta - \mathbf{X}^\top \mathbf{y} = 0$$

$$(\mathbf{X}^\top \mathbf{X})\theta = \mathbf{X}^\top \mathbf{y}$$

Closed Form Solution:

$$\theta = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Closed Form Solution

- Can obtain θ by simply plugging X and y into

$$\theta = (X^T X)^{-1} X^T y$$

$$X = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_d^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & \dots & x_d^{(i)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(n)} & \dots & x_d^{(n)} \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

- If $X^T X$ is not invertible (i.e., singular), may need to:
 - Use pseudo-inverse instead of the inverse
 - In python, `numpy.linalg.pinv(a)`
 - Remove redundant (not linearly independent) features
 - Remove extra features to ensure that $d \leq n$

Gradient Descent vs Closed Form



Gradient Descent

- Requires multiple iterations
- Need to choose α
- Works well when n is large
- Can support incremental learning

Closed Form Solution

- Non-iterative
- No need for α
- Slow if n is large
 - Computing $(X^T X)^{-1}$ is roughly $O(n^3)$

Linear Basis Function Models

Extending Linear Regression to More Complex Models



- The inputs \mathbf{X} for linear regression can be:
 - Original quantitative inputs
 - Transformation of quantitative inputs
 - e.g. log, exp, square root, square, etc.
 - Polynomial transformation
 - example: $y = \beta_0 + \beta_1 \cdot x + \beta_2 \cdot x^2 + \beta_3 \cdot x^3$
 - Basis expansions
 - Dummy coding of categorical inputs
 - Interactions between variables
 - example: $x_3 = x_1 \cdot x_2$

This allows use of linear regression techniques to fit non-linear datasets.

Linear Basis Function Models

- Basic Linear Model:
$$h_{\theta}(\mathbf{x}) = \sum_{j=0}^d \theta_j x_j$$
- Generalized Linear Model:
$$h_{\theta}(\mathbf{x}) = \sum_{j=0}^d \theta_j \phi_j(\mathbf{x})$$
- Once we have replaced the data by the outputs of the basis functions, fitting the generalized model is exactly the same problem as fitting the basic model
 - Unless we use the kernel trick – more on that when we cover support vector machines
 - Therefore, there is no point in cluttering the math with basis functions

Linear Basis Function Models

- Generally,

$$h_{\theta}(\mathbf{x}) = \sum_{j=0}^d \theta_j \underbrace{\phi_j(\mathbf{x})}_{\text{basis function}}$$

- Typically, $\phi_0(\mathbf{x}) = 1$ so that θ_0 acts as a bias
- In the simplest case, we use linear basis functions :

$$\phi_j(\mathbf{x}) = x_j$$

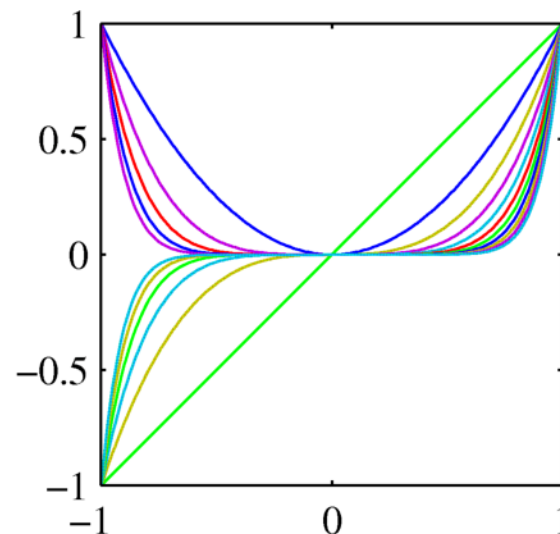
Linear Basis Function Models



- Polynomial basis functions:

$$\phi_j(x) = x^j$$

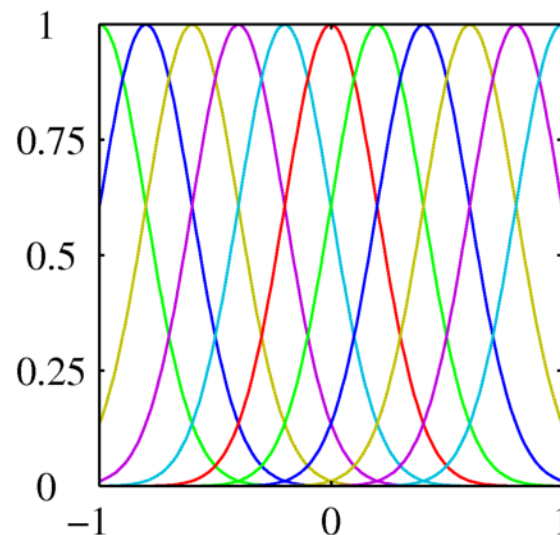
- These are global; a small change in x affects all basis functions



- Gaussian basis functions:

$$\phi_j(x) = \exp \left\{ -\frac{(x - \mu_j)^2}{2s^2} \right\}$$

- These are local; a small change in x only affect nearby basis functions. μ_j and s control location and scale (width).



Linear Basis Function Models




- Sigmoidal basis functions:

$$\phi_j(x) = \sigma \left(\frac{x - \mu_j}{s} \right)$$

where

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- These are also local; a small change in x only affects nearby basis functions. μ_j and s control location and scale (slope).



By using nonlinear basis functions, we allow the function $y(\mathbf{x}, \mathbf{w})$ to be a nonlinear function of the input vector \mathbf{x} . They are called linear models because this function is linear in \mathbf{w} .

Evaluation

- Accuracy
- Precision and recall
- Squared error
- Likelihood
- Posterior probability
- Cost / Utility
- Margin
- Entropy
- etc.

Evaluating Performance

- If y is discrete:
 - Accuracy: $\# \text{ correctly classified} / \# \text{ all test examples}$
 - Good for class balanced dataset
- Want evaluation metric to be in some range, e.g. $[0 \ 1]$
 - 0 = worst possible classifier, 1 = best possible classifier

Confusion matrix



		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

		Predicted: NO	Predicted: YES	
n=165	Actual: NO	TN = 50	FP = 10	60
	Actual: YES	FN = 5	TP = 100	105
		55	110	

Thanks