



BITS Pilani
Pilani Campus

Machine Learning DSECL ZG565

- Dr. Monali Mavani

Topics to be covered



- Instance based learning
- K-Nearest Neighbour Learning
- Locally Weighted Regression (LWR) Learning

Tom Mitchell – Ch 8

Model-based learning techniques

Use the input data

$$\begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n} \end{bmatrix} \text{ and } \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_m \end{bmatrix}$$



To learn a set of parameters

$$\begin{bmatrix} \theta_0 & \theta_1 & \dots & \theta_n \end{bmatrix}$$



Which yield a **generalized** function

$$f(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$



Capable of predicting values or classes on new input data

$$f(x_i) = 39$$

$$f(x_j) = 1$$

Instance-based learning techniques

Store the input data

$$\begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n} \end{bmatrix} \text{ and } \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_m \end{bmatrix}$$



When asked to predict a new value (a query)

$$y_i = ?$$



Search for similar data points previously stored

$$\begin{bmatrix} x_{4,1} & x_{4,2} & \dots & x_{4,n} \\ x_{9,1} & x_{9,1} & \dots & x_{9,n} \\ x_{15,1} & x_{15,1} & \dots & x_{15,n} \end{bmatrix} \text{ and } \begin{bmatrix} y_4 \\ y_9 \\ y_{15} \end{bmatrix}$$



And use them to generate your prediction

$$y_i = \frac{y_4 + y_9 + y_{15}}{3}$$

Instance based learning

- Can approximate real-valued or discrete-valued target functions.
- 'lazy learning', as learning is postponed until a new instance is encountered
- Constructs a local approximation to the target function, applicable in the neighbourhood of new instance
- **Suitable in cases where target function is complex over the entire input space, but easily describable in local approximations**
- A key advantage of this kind of delayed, or lazy, learning is that instead of estimating the target function once for the entire instance space, these methods can estimate it **locally and differently** for each new instance to be classified. E.g Nearest Neighbour and locally weighted regression
- Real world applications found in recommendation systems (amazon).

Eager: generalize before seeing query

- Radial basis function networks, decision trees, back-propagation
- Eager learner must create global approximation



Disadvantages-Instance based approaches

- High cost of classifying new instances
 - nearly all computation takes place at classification time rather than when the training examples are first encountered.
 - Need techniques for efficiently indexing training examples to reduce computation required at query time.

Nearest neighbor approaches

- Key idea: just store all training examples $\langle x_i, f(x_i) \rangle$
- Given query instance x_q , first locate nearest training example x_n , then estimate $f^*(x_q) = f(x_n)$

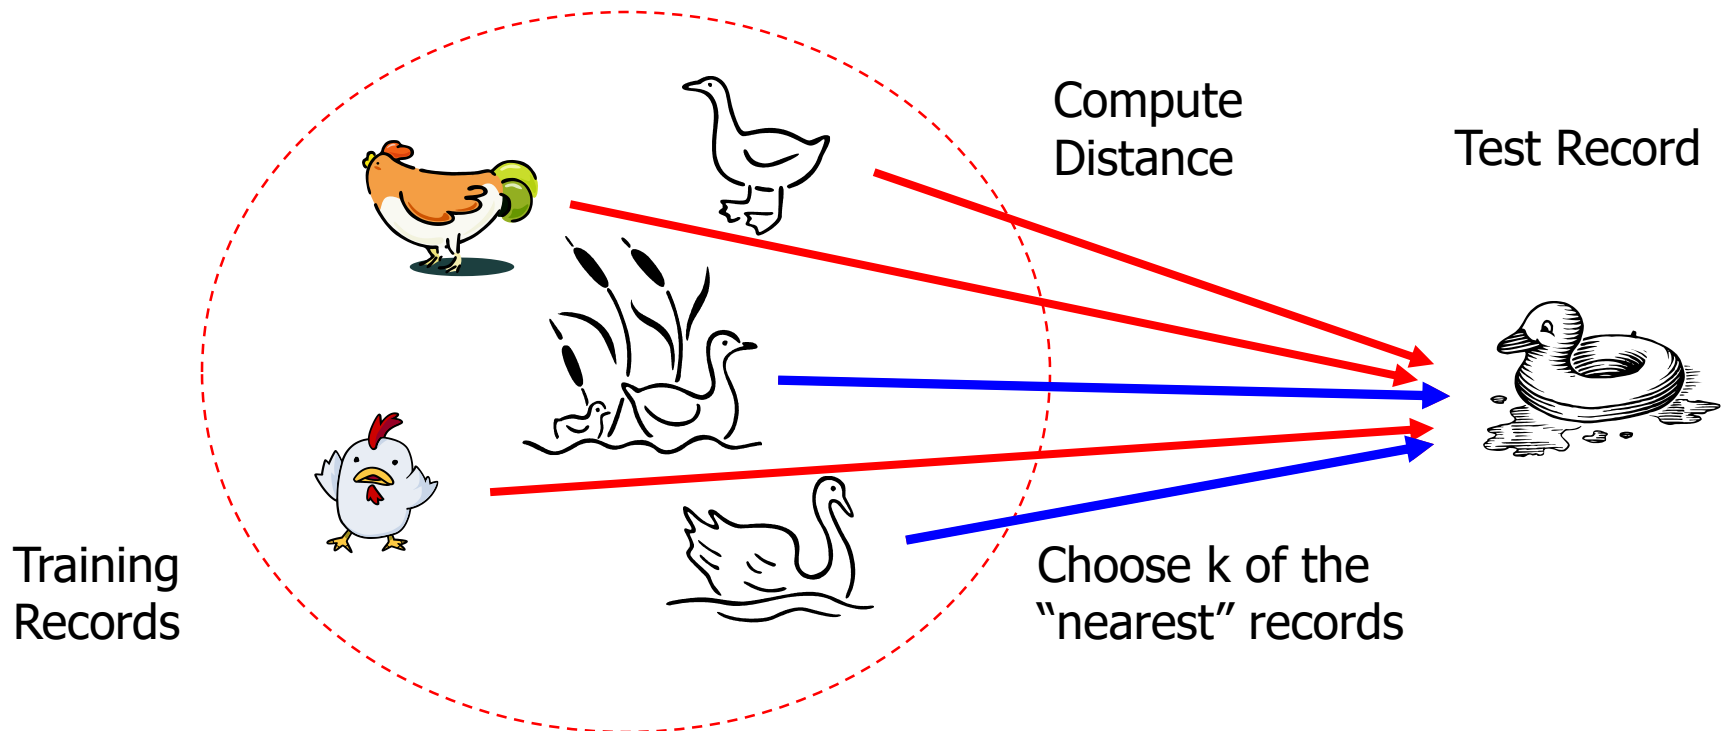
K-nearest neighbor:

- Given x_q , take **vote** among its k nearest neighbors (if **discrete-valued target function**)
- Take **mean(or median)** of f values of k nearest neighbors (if **real-valued**) $f^*(x_q) = \sum_{i=1}^k f(x_i) / k$

Nearest Neighbor Classifiers

Basic idea:

- If it walks like a duck, quacks like a duck, then it's probably a duck



k-Nearest Neighbor

- Considers all instances as members of n-dimensional space
- **Nearest neighbours** of an instance is determined based on Euclidean distance
- Distance between two n-dimensional instances x_i and x_j is given by:


$$d(x_i, x_j) \equiv \sqrt{\sum_{r=1}^n (a_r(x_i) - a_r(x_j))^2}$$

Distances



- Manhattan Distance


$$|X1-X2| + |Y1-Y2|$$



Manhattan

- Euclidean Distance

$$\sqrt{(x1 - x2)^2 + (y1 - y2)^2}$$



Euclidean

- works if the points are arranged in the form of a grid
- if the input variables are not similar in type (such as age, gender, height, etc.)

- Euclidean is commonly used on dense, continuous variables.
- There every dimension matters, and a 20 dimensional space can be challenging
- if the input variables are of similar in type (such as width, height, depth etc.)

Distance Measures : Special Cases of Minkowski

- $h = 1$: **Manhattan** (city block, L_1 norm) **distance**

$$d(i, j) = |x_{i_1} - x_{j_1}| + |x_{i_2} - x_{j_2}| + \dots + |x_{i_p} - x_{j_p}|$$

- $h = 2$: (L_2 norm) **Euclidean** distance

$$d(i, j) = \sqrt{(|x_{i_1} - x_{j_1}|^2 + |x_{i_2} - x_{j_2}|^2 + \dots + |x_{i_p} - x_{j_p}|^2)}$$

- $h \rightarrow \infty$. **“supremum”** (L_{\max} norm, L_{∞} norm) distance, **Chebyshev distance**.
 - This is the maximum difference between any component (attribute) of the vectors

$$d(i, j) = \lim_{h \rightarrow \infty} \left(\sum_{f=1}^p |x_{if} - x_{jf}|^h \right)^{\frac{1}{h}} = \max_f^p |x_{if} - x_{jf}|$$

- Different features may have different measurement scales
 - E.g., patient weight in kg (range [50,200]) vs. blood protein values in ng/dL (range [-3,3])
- Consequences
 - Patient weight will have a much greater influence on the distance between samples
 - May bias the performance of the classifier

Feature scaling - - Standardization



- X : raw score to be standardized, μ : mean of the population, σ : standard deviation
- the distance between the raw score and the population mean in units of the standard deviation
- negative when the raw score is below the mean, “+” when above Where

$$z = \frac{x - \mu}{\sigma}$$

- The Standard Scaler assumes data is normally distributed within each feature and scales them such that the distribution centered around 0, with a standard deviation of 1.

Feature scaling - Normalization



Min-Max scaler/Normalization

Feature scaling of features x_i consists of rescaling the range of features to scale the range in $[0, 1]$ or $[-1, 1]$ (Do not apply to x_0)

$$x_1 = \frac{\text{size} - 1000}{2000}$$

Average value of x_1

Maximum value of x_1 – min value of x_1

$$x_2 = \frac{\#bedrooms - 2}{5}$$

Discrete and Continuous-valued function

● discrete-valued target function:

- $f : \mathcal{R}^n \rightarrow V$ where V is the finite set $\{v_1, v_2, \dots, v_s\}$
- the target function value is the most common value among the k nearest training examples

$$\hat{f}(x_q) \leftarrow \underset{v \in V}{argmax} \sum_{i=1}^k \delta(v, f(x_i))$$

where $\delta(a, b) = (a == b)$

● continuous-valued target function:

- algorithm has to calculate the mean value instead of the most common value
- $f : \mathcal{R}^n \rightarrow \mathcal{R}$

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k f(x_i)}{k}$$

k-Nearest Neighbor Classifier



Training algorithm:

- For each training example $\langle x, f(x) \rangle$, add the example to the list *training_examples*

Classification algorithm:

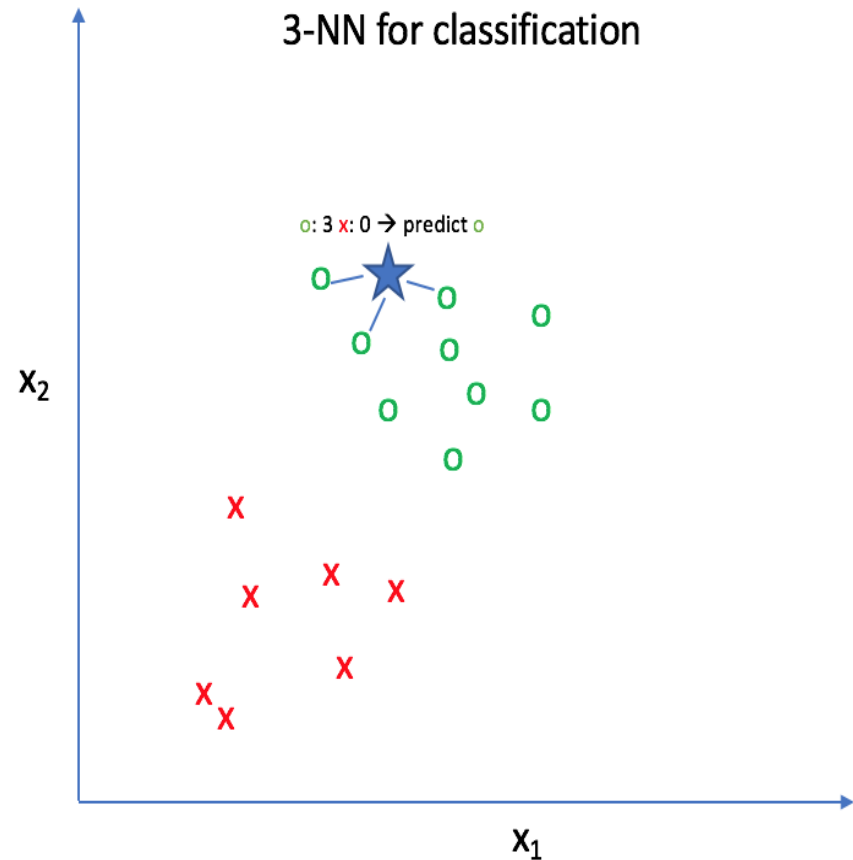

- Given a query instance x_q to be classified,
 - Let $x_1 \dots x_k$ denote the k instances from *training_examples* that are nearest to x_q
 - Return

$$\hat{f}(x_q) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k \delta(v, f(x_i))$$


where $\delta(a, b) = 1$ if $a = b$ and where $\delta(a, b) = 0$ otherwise.

* It can be used for Regression as well.


KNN for regression and Classification



k-NN examples



K=1




K=5

Voronoi Diagram

query point q_f

nearest neighbor q_i

- Nearest neighbour approach induces a **Voronoi tessellation**/partition of the input space (all test points falling in a cell will get the label of the training input in that cell)
- For any sample, the nearest sample is determined by the closest Voronoi cell edge




Various issues that affect the performance of kNN:



Performance of a classifier largely depends on the of the hyperparameter k

- Choosing smaller values for K , noise can have a higher influence on the result.
- Larger values of k are computationally expensive

Assigning the class labels can be tricky. For example, in the below case, for ($k=5$) the point is closer to 'green' classification, but gets classified as 'red' due to higher red votes/majority voting to 'red'




Value of K

- Choosing the value of k:
 - If k is too small, sensitive to noise points
 - If k is too large, neighborhood may include points from other classes

Rule of thumb:


$K = \sqrt{N}$

N: number of training points



Finding K - Elbow method

- Compute sum of squares error (SSE) or any other error function for varying values of K (1 to a reasonable X) and plot against K
- In the plot, the elbow (see pic) gives the value of K beyond which the error function plot almost flattens
- As K approaches the total number of instances in the set, error function drops down to '0',



Distance weighted nearest neighbor

- contribution of each of the k nearest neighbors is weighted accorded to their distance to x_q
 - **discrete-valued target functions**

$$\hat{f}(x_q) \leftarrow \underset{v \in V}{argmax} \sum_{i=1}^k w_i \delta(v, f(x_i))$$

where $w_i \equiv \frac{1}{d(x_q, x_i)^2}$ and $\hat{f}(x_q) = f(x_i)$ if $x_q = x_i$

- **continuous-valued target function:**

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k w_i f(x_i)}{\sum_{i=1}^k w_i}$$

Distance Weighted k-NN

- Give more weight to neighbors closer to the query point
 - $w_i = K(d(x_q, x_i))$
 - $K(d(x_q, x_i)) = 1 / d(x_q, x_i)^2$
 - $K(d(x_q, x_i)) = 1 / (d_0 + d(x_q, x_i))^2$
 - $K(d(x_q, x_i)) = \exp(-(d(x_q, x_i) / \sigma_0)^2)$
 - $d(x_q, x_i)$ is the distance between x_q and x_i
- Variation: Instead of only k-nearest neighbors use all training examples (global method - Shepard's method)

Distance Weighted Average



Weighting the error criterion

$$-E(x_q) = \sum_i (f^*(x_q) - f(x_i))^2 K(d(x_i, x_q))$$

Best estimate $f^*(x_q)$ will minimize the cost $E(x_q)$,
therefore $\partial E(x_q) / \partial f^*(x_q) = 0$

Curse of Dimensionality



- Imagine instances described by 20 attributes but only a few(2) are relevant to target function
- *Curse of dimensionality*: nearest neighbor is easily misled when instance space is high-dimensional
- First approach:
 - completely eliminate the least relevant attributes from the instance space. This is equivalent to setting some of the z_i scaling factors to zero
 - Use of cross-validation methods for selecting relevant subsets of the attributes

Second approach



- weight each attribute differently when calculating the distance between two instances i.e.
 - stretching the axes in the Euclidean space, shortening the axes that correspond to less relevant attributes, and lengthening the axes that correspond to more relevant attributes.
- The amount by which each axis should be stretched can be determined automatically using a cross-validation approach
- To stretch (multiply) the j th axis by some factor z_j , where the values $z_1 \dots z_n$, are chosen to minimize the true classification error of the learning algorithm

When to Consider Nearest Neighbors



- Suitable for Low dimensional datasets
- Lots of training data (distance-weighted KNN)
 - Training is very fast
- Learn complex target functions
 - Do not lose information
- Noisy training data (distance-weighted KNN)
 - by taking the weighted average of the k neighbors nearest to the query point, it can smooth out the impact of isolated noisy training examples.

Nearest-Neighbor Classifiers: Issues

- The value of k , the number of nearest neighbors to retrieve
- Choice of Distance Metric to compute distance between records
- They typically consider ***all*** attributes of the instances when attempting to retrieve similar training examples from memory.
 - If the target concept depends on only a few of the many available attributes, then the instances that are truly most "similar" may well be a large distance apart.



Nearest Neighbours issues

- Expensive, Slow at query time
 - To determine the nearest neighbour of a query point q , must compute the distance to all N training examples
 - + Pre-sort training examples into fast data structures (kd-trees)
 - + Remove redundant data (condensing)
- Storage Requirements
 - Must store all training data **P**
 - + Remove redundant data (condensing)
 - Pre-sorting often increases the storage requirements
- High Dimensional Data
 - “Curse of Dimensionality”
 - Required amount of training data increases exponentially with dimension
 - Computational cost also increases dramatically

Locally Weighted Regression

Locally Weighted Regression



- Locally – Function approximated based on data near query point
- Weighted – Contribution by each training example is weighted by its distance from query point
- Regression- Approximates real-valued target function
- **Residual** is the error in approximating the target function.

$$\hat{f}(x) - f(x)$$

- **Kernel function** is the function of distance that is used to determine the weight of each training example. In other words, the kernel function is the function **K** such that


$$w_i = K(d(x_i, x_q)).$$




Locally Weighted Regression

- The nearest-neighbor approaches described in the previous section can be thought of as approximating the target function $f(\mathbf{x})$ at the single query point $\mathbf{x} = \mathbf{x}_q$.
- Locally weighted regression is a generalization of this approach. It constructs an explicit approximation to f over a local region surrounding \mathbf{x}_q .
- Uses nearby or distance-weighted training examples to form this local approximation to f .
- Approximate the target function in the neighborhood surrounding \mathbf{x} , using a linear function, a quadratic function, a multilayer neural network, or some other functional form.
 - More complex functional forms are not often found the cost of fitting more complex functions for each query instance is prohibitively high,
 - these simple approximations model the target function quite well over a sufficiently small subregion of the instance space

Example



3-nearest-neighbors linear regression



locally weighted regression

Locally weighted linear regression



- target function is approximated using a **linear function**

$$\hat{f}(x) = w_0 + w_1 a_1(x) + \dots + w_n a_n(x)$$

- methods like **gradient descent** can be used to calculate the coefficients w_0, w_1, \dots, w_n to minimize the error in fitting such linear functions
 - ANNs require a global approximation to the target function
 - here, just a local approximation is needed
- ⇒ the error function has to be redefined

Locally weighted linear regression



• possibilities to redefine the error criterion E

1. Minimize the squared error over just the k nearest neighbors

$$E_1(x_q) \equiv \frac{1}{2} \sum_{x \in k \text{ nearest neighbors}} (f(x) - \hat{f}(x))^2$$

2. Minimize the squared error over the entire set D , while weighting the error of each training example by some decreasing function K of its distance from x_q

$$E_2(x_q) \equiv \frac{1}{2} \sum_{x \in D} (f(x) - \hat{f}(x))^2 \cdot K(d(x_q, x))$$

3. Combine 1 and 2

$$E_3(x_q) \equiv \frac{1}{2} \sum_{x \in k \text{ nearest neighbors}} (f(x) - \hat{f}(x))^2 \cdot K(d(x_q, x))$$

Locally weighted linear regression



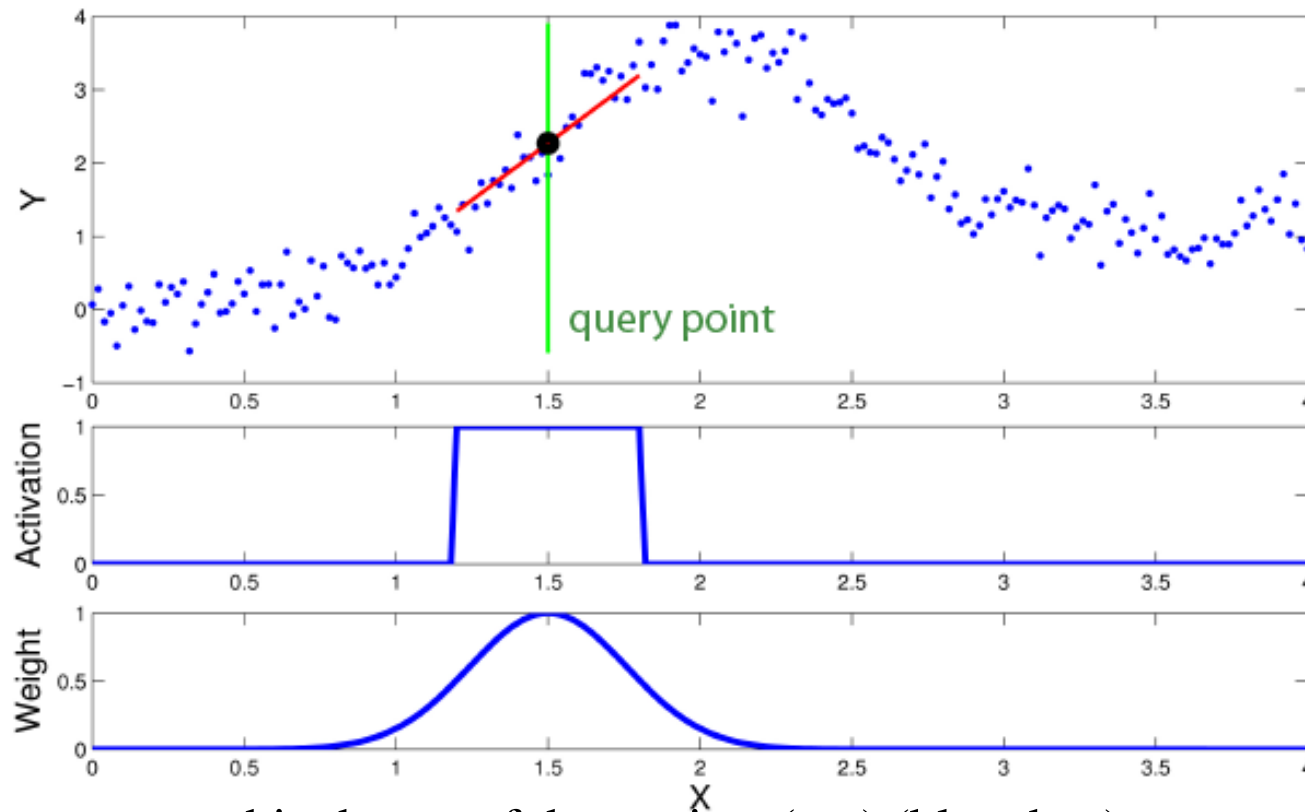
- For a given query point \mathbf{x}_q we solve the following weighted regression problem using gradient descent:

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_j \mathcal{K}(\operatorname{Distance}(\mathbf{x}_q, \mathbf{x}_j)) (y_j - \mathbf{w} \cdot \mathbf{x}_j)^2$$

$$h(\mathbf{x}_q) = \mathbf{w}^* \cdot \mathbf{x}_q$$

- Note that we need to solve a new regression problem for *every* query point—that's what it means to be *local*.
- In ordinary linear regression, we solved the regression problem once, globally, and then used the same $h_{\mathbf{w}}$ for any query point.

Example



- The upper graphic the set of data points (x,y) (blue dots), query point (green line), local linear model (red line) and prediction (black dot).
- The graphic in the middle shows the activation area of the model.
- The corresponding weighting kernel (receptive field) is shown in the bottom graphic.

Thank You