



BITS Pilani
Pilani | Dubai | Goa | Hyderabad

DSECL ZG 522: Big Data Systems

Session 12: Storage on Cloud (AWS examples)


Janardhanan PS

Professor

janardhanan.ps@wilp.bits-pilani.ac.in



Amazon EBS
(persistent)



Amazon EC2
Instance Store
(ephemeral)

Block



Amazon EFS

File



Amazon S3



Amazon Glacier

Object

storage with
meta-data

Data Transfer



AWS
Snow Family



AWS Storage
Gateway



EFS
File Sync



3rd Party
Connectors



AWS Direct
Connect



S3 Transfer
Acceleration



Amazon
Kinesis

Topics for today

- **Object storage**
 - ✓ S3
- **File storage**
 - ✓ Elastic File System (EFS)
- **Block storage**
 - ✓ Elastic Block Storage (EBS)
- **Databases**
 - ✓ Key-value: DynamoDB



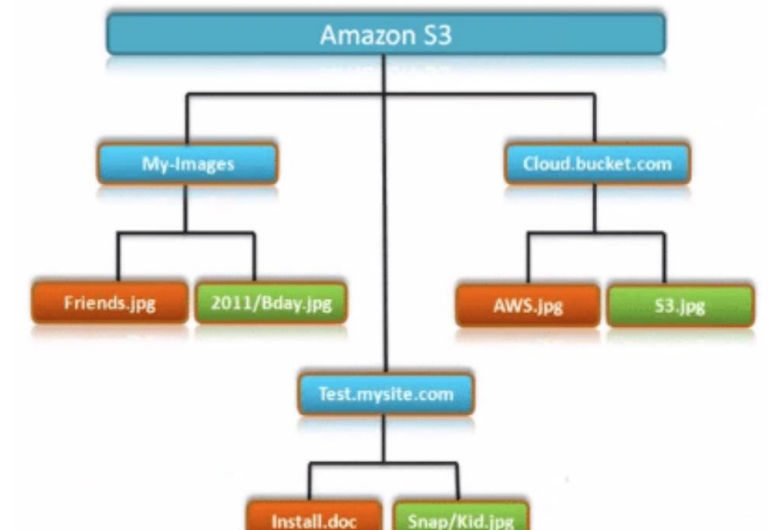
S3 - Object storage

- S3 is an object storage (Simple Storage Service)
 - ✓ which means it stores objects which are files with lot of meta-data associated
- A bucket is the container for objects
- An S3 account can have hundreds of buckets and a bucket can have hundreds of objects
- A bucket can also have folders to organise objects
- An object
 - ✓ can be 1 byte to 5TB
 - ✓ is uniquely identified by a developer assigned key and a URL
 - ✓ has an ACL to control who can access from anywhere - not necessarily from within AWS
 - ✓ supports versioning and “eventual consistency” across multiple reads / writes
 - ✓ is partitioned and replicated

Example

- 3 Buckets with folders inside buckets for grouping objects
- Can be access via public URL:
 - ✓ Option 1:
bucketname.s3.amazonaws.com/objectname
 - ✓ Option 2:
s3.amazonaws.com/bucketname/objectname

Amazon S3 Namespace




Consistency model

- S3 - Follows AP design wrt CAP Theorem
- So when a user uploads an object, it is replicated but readers can access inconsistent replicas and all replicas are “eventually consistent”
- Option for reduced redundancy for lower cost - so writes / updates may not be durable but cost is reduced significantly
 - ✓ RRS: <https://aws.amazon.com/blogs/aws/new-amazon-s3-reduced-redundancy-storage-rrs/>
- Some observations:
 - ✓ User uploads object but a reader can get “key does not exist”
 - ✓ An object may not appear in a listing of a bucket immediately
 - ✓ Old data may be returned if a write is not propagated by then
 - ✓ A deleted object may be seen in a listing for some time

Architecture

- 20+ Regions with multiple Availability Zones (AZ) per region
 - ✓ AZs are physically isolated connected over low latency network
- 60+ AZ with each AZ is upto 8 DCs
- Data can be stored in at least 3 physically separated AZ within a Region for HA
- Private network connections across AZ and DCs for low latency



Features

- Storage tiering
- Object Lambda
- Access points associated with ACLs for customised client access
- Batch mode operations




Storage classes



S3 Standard



S3 Standard –
Infrequent Access



Amazon Glacier

Active data
Synchronous access
Milliseconds retrieval
2.1¢-GB/mo

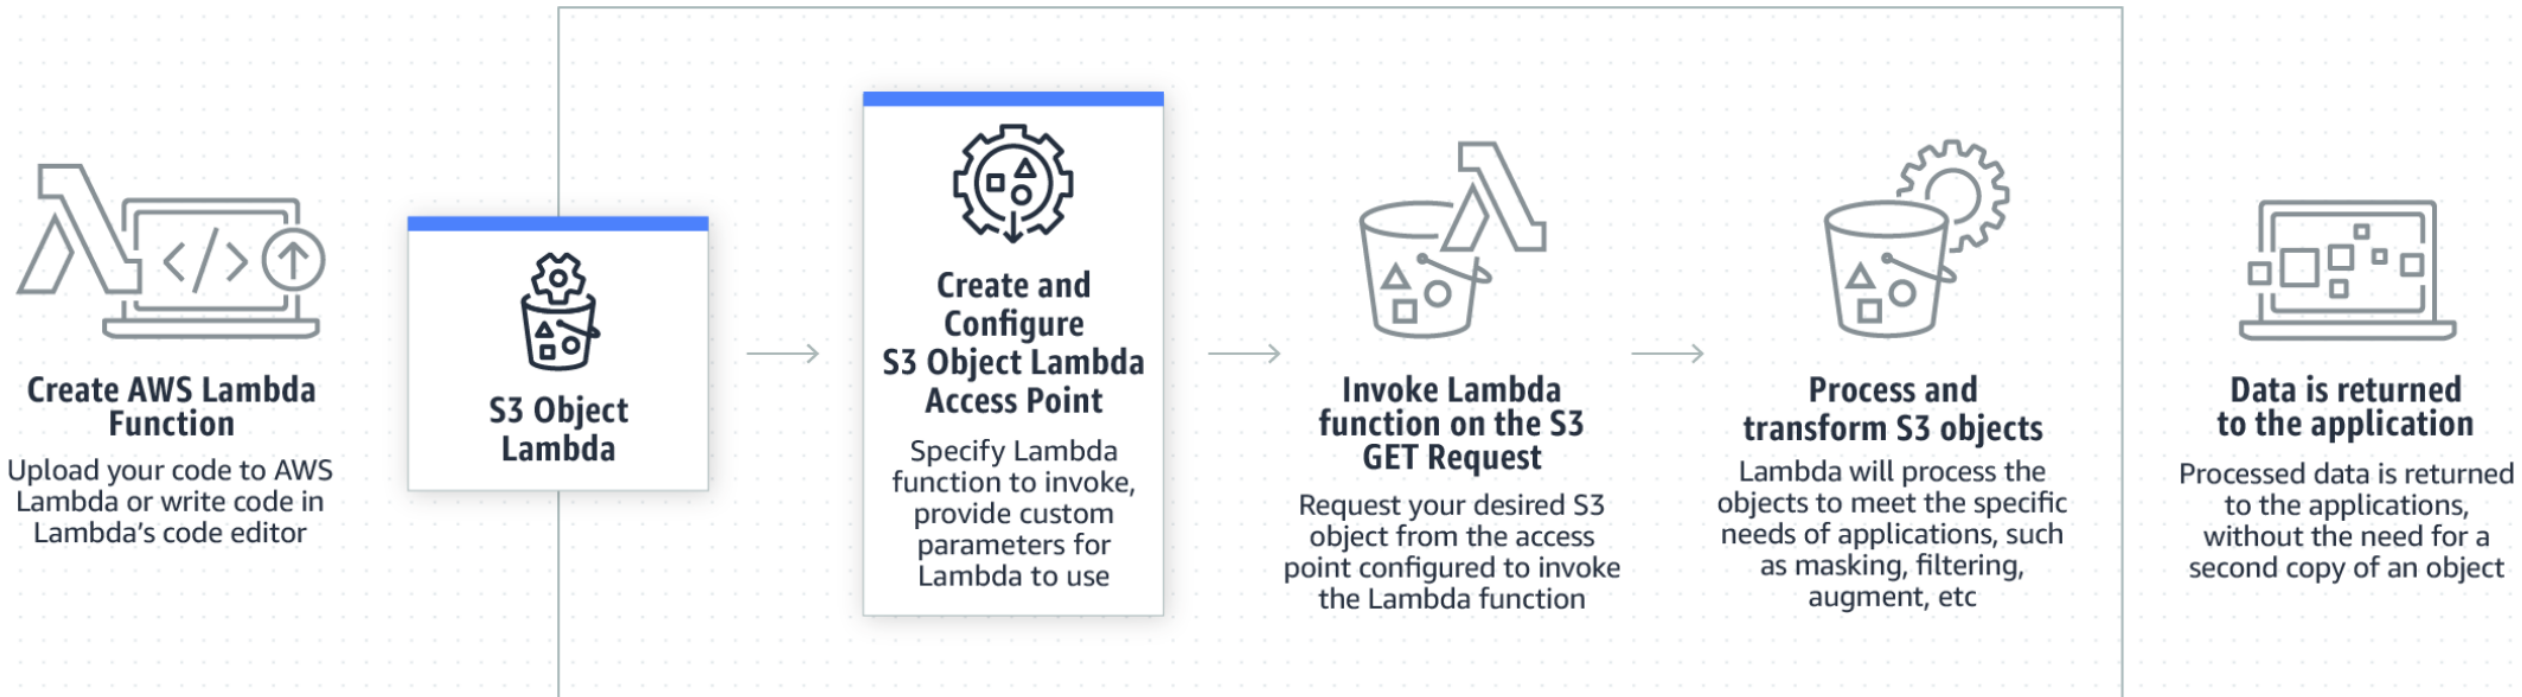
Infrequently accessed data
Synchronous access
Milliseconds retrieval
1.25¢-GB/mo

Archive data
Asynchronous access
Minutes-to-hours retrieval
0.4¢-GB/mo

S3 Outposts also enables to store data on customer premises for faster local access or local residency requirements

Object Lambda

Remember Function-as-a-Service ? What if you could move the function closer to Data ?



Object Lambda

- Add code to S3 GET request for data processing

```
GET /my-image.jpg HTTP/1.1
Host: bucket.s3.<Region>.amazonaws.com
Date: Mon, 3 Oct 2016 22:32:00 GMT
Authorization: authorization string
```

- Register an access point with Lambda transformation function

```
PUT /v20180820/accesspointforobjectlambda/<name> HTTP/1.1
x-amz-account-id: <AccountId>
<?xml version="1.0" encoding="UTF-8"?>
<CreateAccessPointForObjectLambdaRequest xmlns="http://awss3control.amazonaws.com/doc/2018-08-20/">
  <Configuration>
    <AllowedFeatures>
      <AllowedFeature><string></AllowedFeature>
    </AllowedFeatures>
    <CloudWatchMetricsEnabled><boolean></CloudWatchMetricsEnabled>
    <SupportingAccessPoint><string></SupportingAccessPoint>
    <TransformationConfigurations>
      <TransformationConfiguration>
        <Actions>
          <Action><string></Action>
        </Actions>
        <ContentTransformation>
          <AwsLambda>
            <FunctionArn><string></FunctionArn>
            <FunctionPayload><string></FunctionPayload>
          </AwsLambda>
        </ContentTransformation>
      </TransformationConfiguration>
    </TransformationConfigurations>
  </Configuration>
</CreateAccessPointForObjectLambdaRequest>
```



Use cases

- Backup and restore for Cloud as well as on-prem data
- Disaster recovery with cross region replication
- Data archival
- Cloud based applications with scalable storage access from anywhere
- Create data lakes for big data analytics




Topics for today

- Object storage
 - ✓ S3
- **File storage**
 - ✓ Elastic File System (EFS)
- Block storage
 - ✓ Elastic Block Storage (EBS)
- Databases
 - ✓ Key-value: DynamoDB



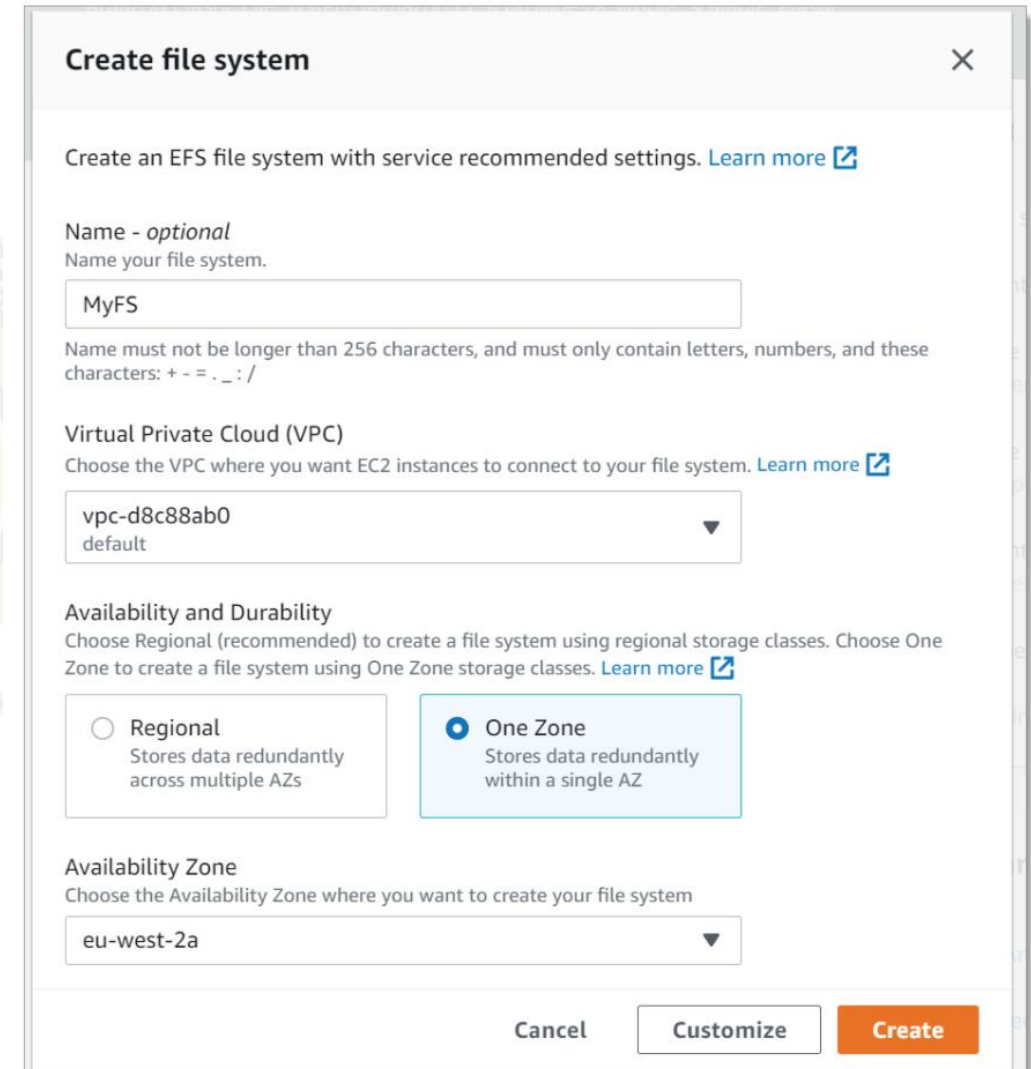
Overview

- File system built using SSD storage that can be accessed from multiple Cloud instances or even customer premises servers
 - ✓ Supports NFS protocol
- Applications can simply attach the file system within VPC - so not much change to get elastic scalability and IOPS
 - ✓ 10GB/sec, 500K IOPS
 - ✓ Multiple NFS clients can attach EFS
- Multiple storage classes



EFS Storage classes

- Standard - with replication to 3+ AZs
 - ✓ EFS Standard and EFS Standard - Infrequent Access
- One Zone - replication only within one AZ at lower cost
 - ✓ EFS One Zone and EFS One Zone - Infrequent Access
 - ✓ 80% files typically in this category
- Setup lifecycle policies to move data based on age to Infrequent Access (IA) class



Create file system

Create an EFS file system with service recommended settings. [Learn more](#)

Name - optional
Name your file system.

Name must not be longer than 256 characters, and must only contain letters, numbers, and these characters: + - = . _ : /

Virtual Private Cloud (VPC)
Choose the VPC where you want EC2 instances to connect to your file system. [Learn more](#)

default

Availability and Durability
Choose Regional (recommended) to create a file system using regional storage classes. Choose One Zone to create a file system using One Zone storage classes. [Learn more](#)

☐ **Regional**
Stores data redundantly across multiple AZs

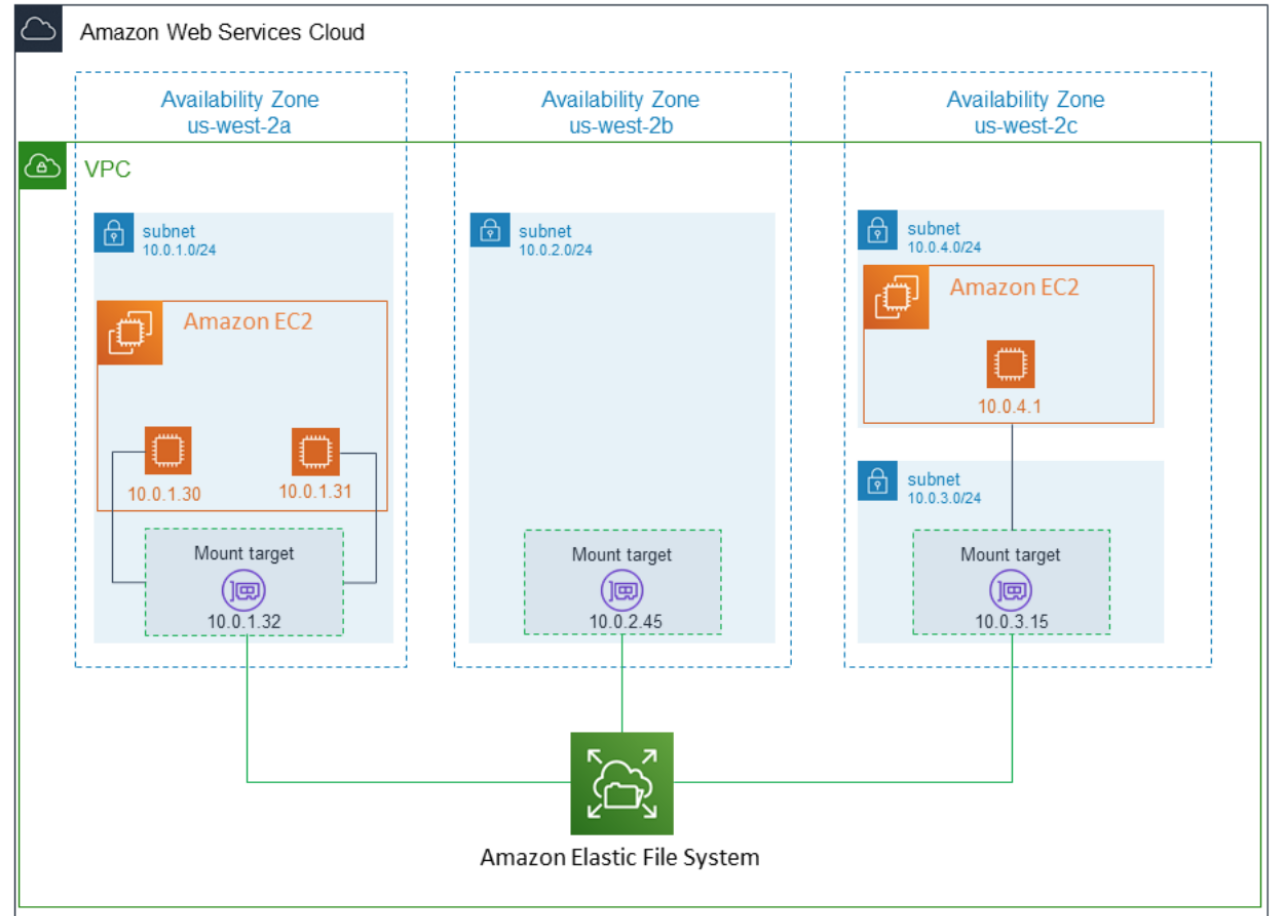
☒ **One Zone**
Stores data redundantly within a single AZ

Availability Zone
Choose the Availability Zone where you want to create your file system

[Cancel](#) [Customize](#) [Create](#)

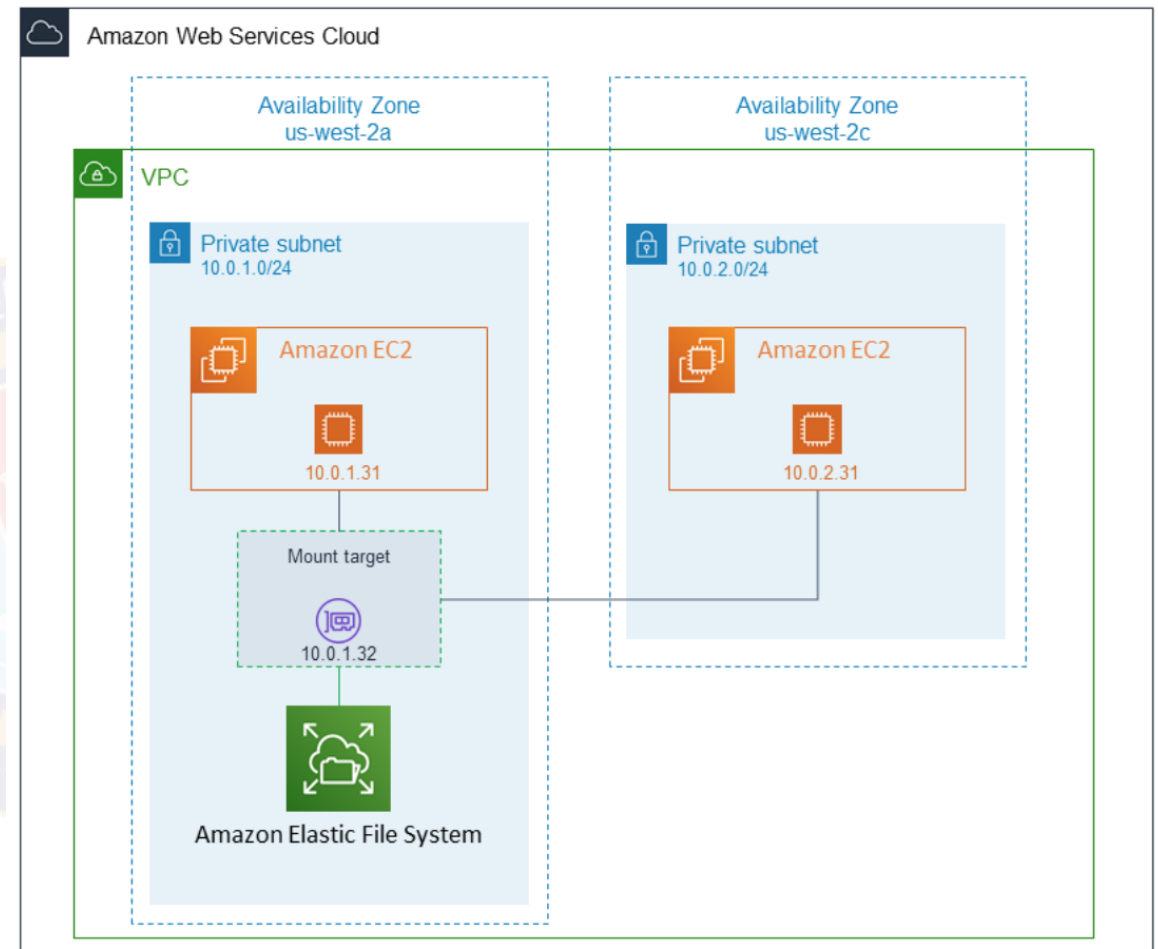
Standard storage class

- A VPC has 3 AZs
- Each AZ has a mount target
- Typically should access EFS from a mount target within same AZ for performance and cost
- Can create a mount target in one of the subnets within an AZ



One Zone storage class

- Single mount target in 1 AZ
- So instance in another AZ has to pay for data access cost



What if you have to connect customer premise ?

Read: <https://docs.aws.amazon.com/efs/latest/ug/how-it-works.html>

Topics for today

- Object storage
 - ✓ S3
- File storage
 - ✓ Elastic File System (EFS)
- **Block storage**
 - ✓ Elastic Block Storage (EBS)
- Databases
 - ✓ Key-value: DynamoDB



Overview

- Block storage for an EC2 instance that can be attached and detached anytime
- Provides a volume as a collection of network attached blocks that are exposed as disks
- Depending on performance / cost one can opt for HDD-backed or SSD-backed volumes
- EBS volumes are durable and replicated within AZ
- Can't move a volume to another AZ without snapshot

EBS or EFS

- EBS

- ✓ Can be only accessed by one instance at a time
- ✓ Steady predictable performance for a single instance use case
- ✓ upto: 4GBps, 64TB, 260K IOPS, sub-millisec latency per volume
- ✓ I/O intensive applications, e.g. relational databases, OLAP engines
- ✓ Multiple performance / cost options
- ✓ Cheaper than EFS per volume but only for one instance - so effectively more expensive

- EFS

- ✓ Like a distributed multi-user network file system
- ✓ Scalable access across many users with decent performance
- ✓ Costs more per GB but is shared by multiple instances - so turns out cheaper for cost sensitive shared storage applications

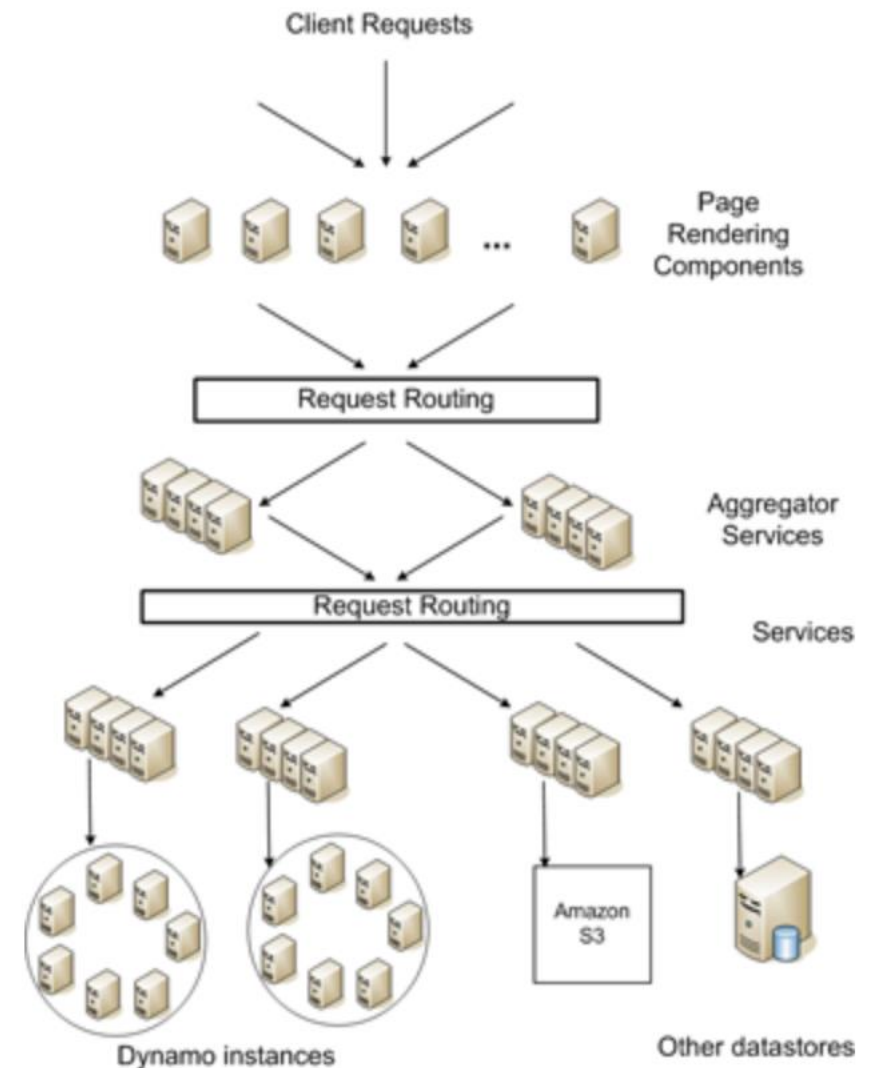
Topics for today

- Object storage
 - ✓ S3
- File storage
 - ✓ Elastic File System (EFS)
- Block storage
 - ✓ Elastic Block Storage (EBS)
- **Databases**
 - ✓ Dynamo paper: https://www.allthingsdistributed.com/2007/10/amazons_dynamo.html
(concepts map to Cassandra, DynamoDB etc.)
 - ✓ DynamoDB:
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>



Applications and requirements

- Best seller lists, shopping carts, customer preferences, session management, sales rank, and product catalog
- Query: no relational queries, simple key based object retrieval with object size < 1 MB
- ACID: Transactions reduce availability. No strict consistency requirement.
- Efficiency: Commodity hardware, strict low latency at 99.9th percentile, tradeoff performance, cost, availability, durability



Design considerations


- Highly available for writes, e.g. update shopping cart even on network partitions
- Complexity of conflict resolution is pushed to read
- Who resolves conflicts on read ?
 - ✓ If data store then only choice is “last write wins” or leave it to the application developer
- Incrementally add nodes and all nodes have symmetric role, like peers
- Heterogenous system - work is proportional to capability

Architectural techniques

| Problem | Technique | Advantage |
|------------------------------------|---|---|
| Partitioning | Consistent Hashing | Incremental Scalability |
| High Availability for writes | Vector clocks with reconciliation during reads | Version size is decoupled from update rates. |
| Handling temporary failures | Sloppy Quorum and hinted handoff | Provides high availability and durability guarantee when some of the replicas are not available. |
| Recovering from permanent failures | Anti-entropy using Merkle trees reduce disorder | Synchronizes divergent replicas in the background. subtree roots keep track of diff of hashes in children |
| Membership and failure detection | Gossip-based membership protocol and failure detection. | Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information. |

Object storage and interface

- Objects stored with associated keys
 - ✓ $\text{object, context} = \text{get}(\text{key})$
 - ✓ $\text{put}(\text{key, context, object})$
 - ✓ context encodes system meta-data along with versions
- Partitioning and replication
 - ✓ A variant of “consistent hashing” of the key is used to allocate storage nodes
 - ✓ Each data item is replicated on a configurable number of nodes




A Chord DHT Example



Legend

- : key = hash (object name)
- : node = hash (IP address)
- : point (potential key or node)

finger table: points to where the object can be found or where the next pointer to the object is

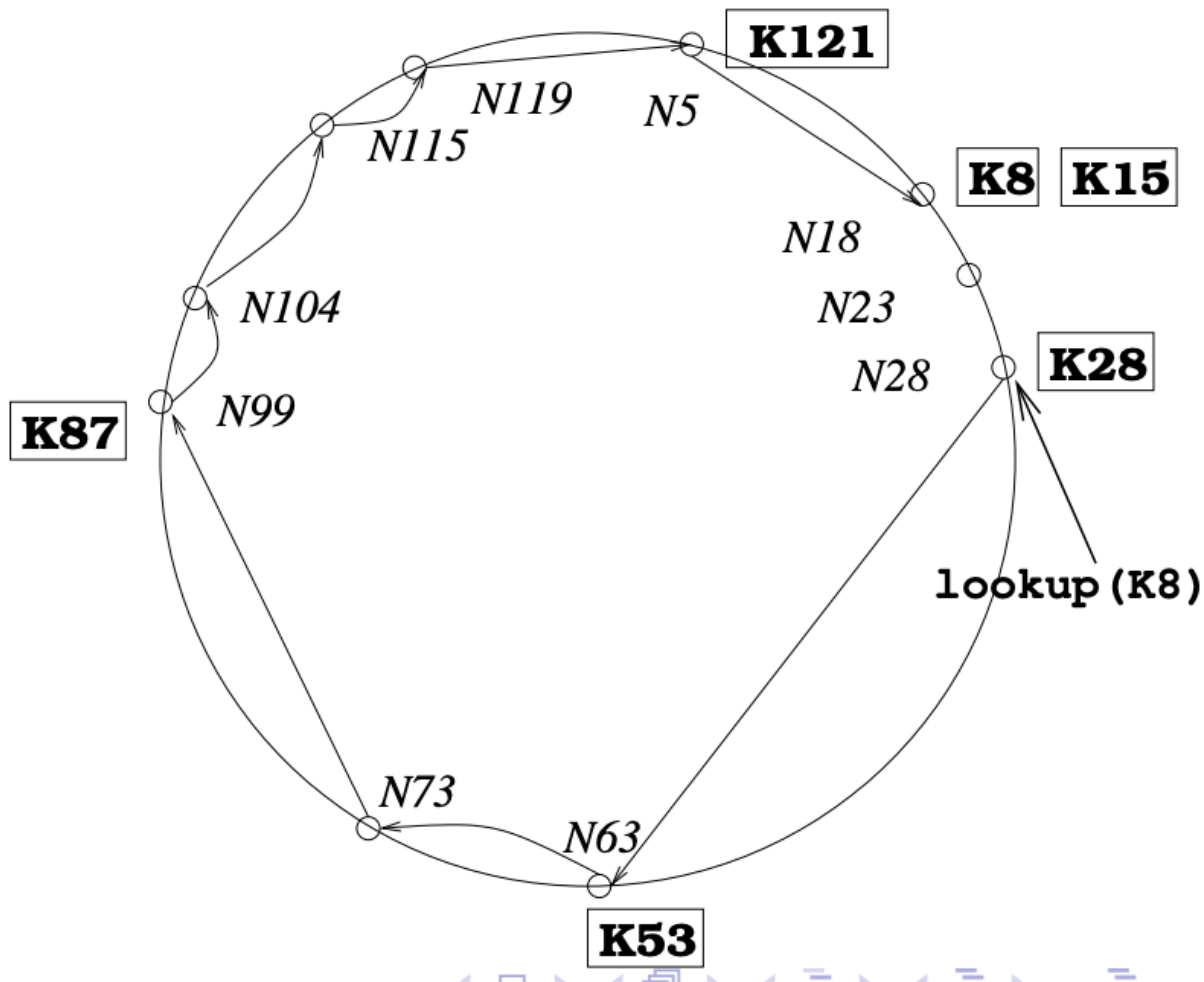
| Key | Reference |
|-----|----------------------|
| 14 | (110.34.56.20, 5200) |
| ... | ... |

 **N17**
80.201.52.40



look for first node equal / following k14: succ(k14)=N17

Linear search example




- Lookup(x) goes to a node
- Then keep following the successor link of nodes to check if object exists
- $O(n)$ message hops
- Routing table in each node is $O(1)$

Chord Finger table for log search

| i | Target Key | Successor of Target Key | Information about Successor |
|----------|---------------|----------------------------|----------------------------------|
| 1 | $N + 1$ | Successor of $N + 1$ | IP address and port of successor |
| 2 | $N + 2$ | Successor of $N + 2$ | IP address and port of successor |
| \vdots | \vdots | \vdots | \vdots |
| m | $N + 2^{m-1}$ | Successor of $N + 2^{m-1}$ | IP address and port of successor |

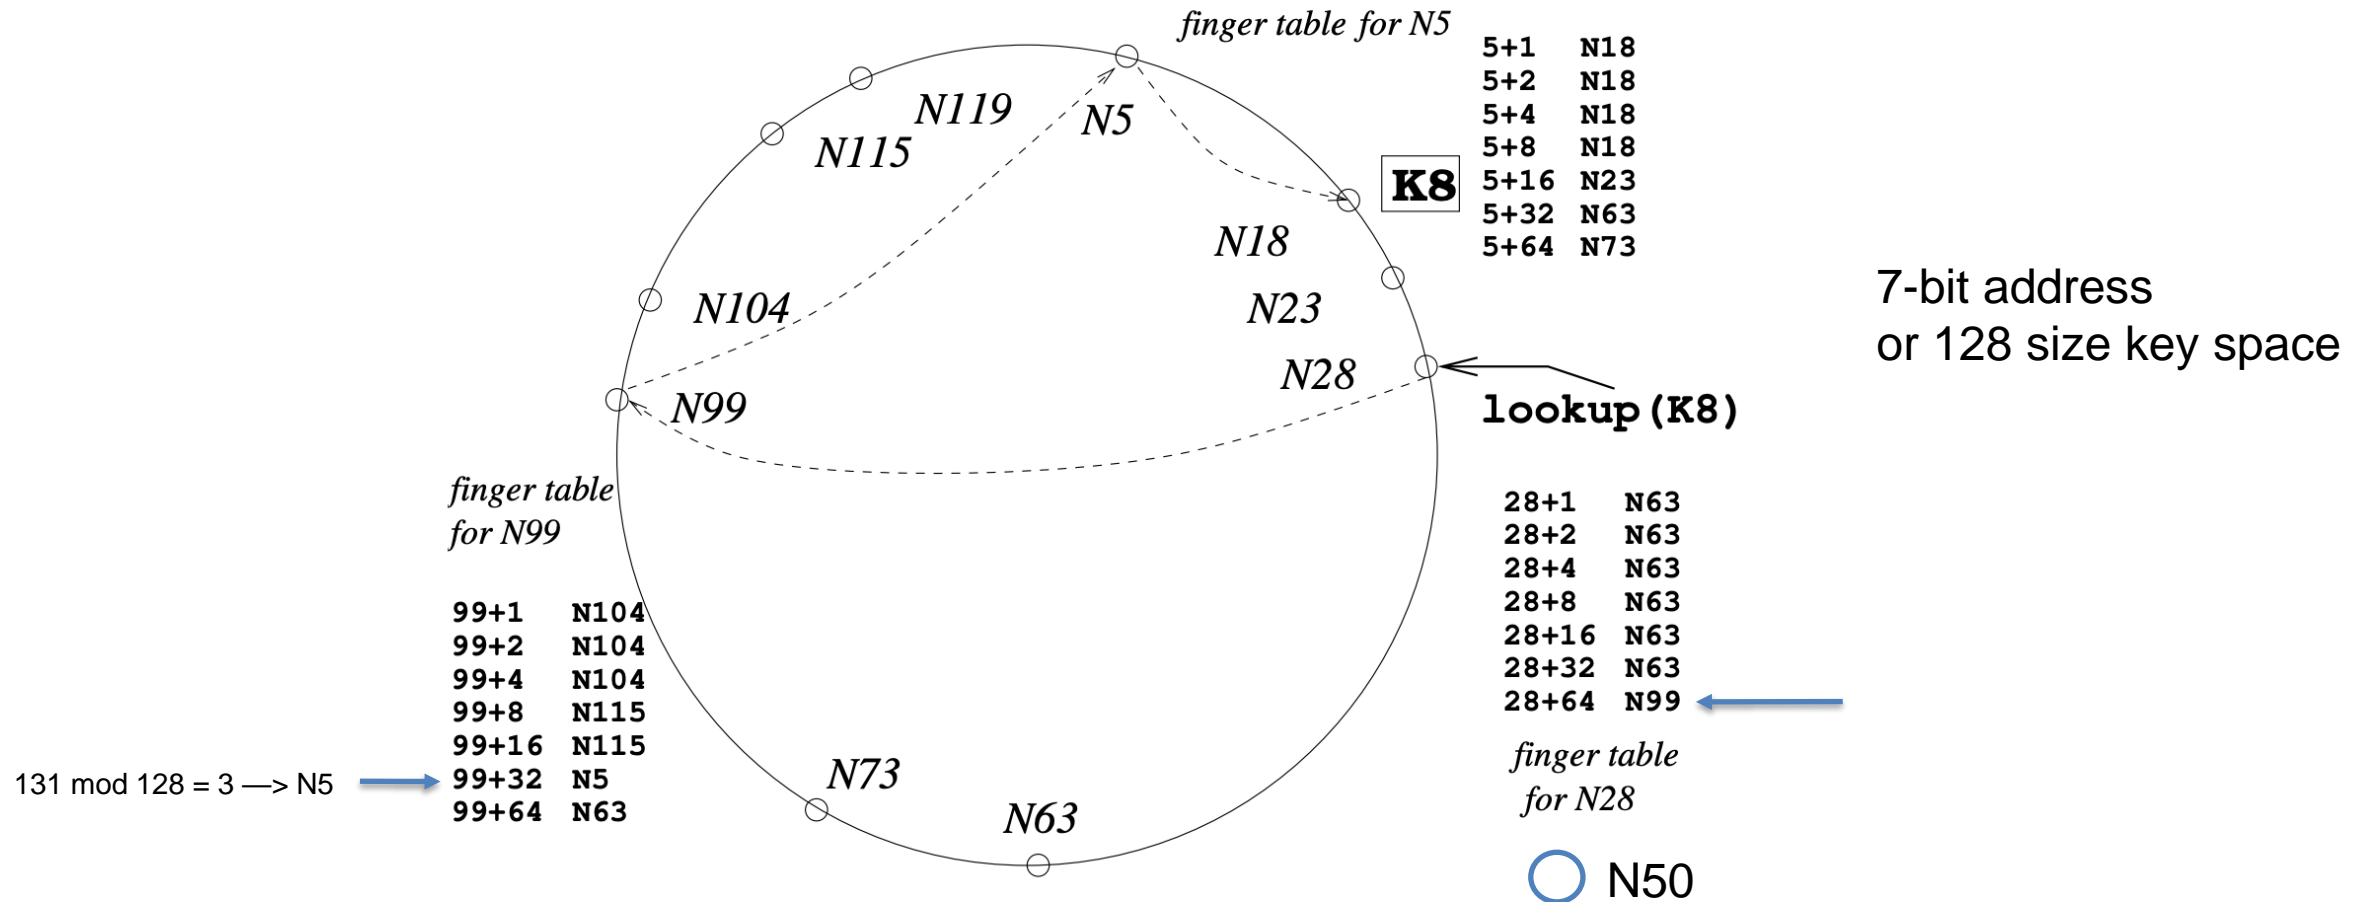
- Each node stores a routing table which helps to search with $O(\log n)$ message hops
- Space taken in routing table is $O(m)$
- When a node joins or leaves, the finger tables in some nodes need to be updated so that the change doesn't impact



https://pdos.csail.mit.edu/papers/chord:sigcomm01/chord_sigcomm.pdf

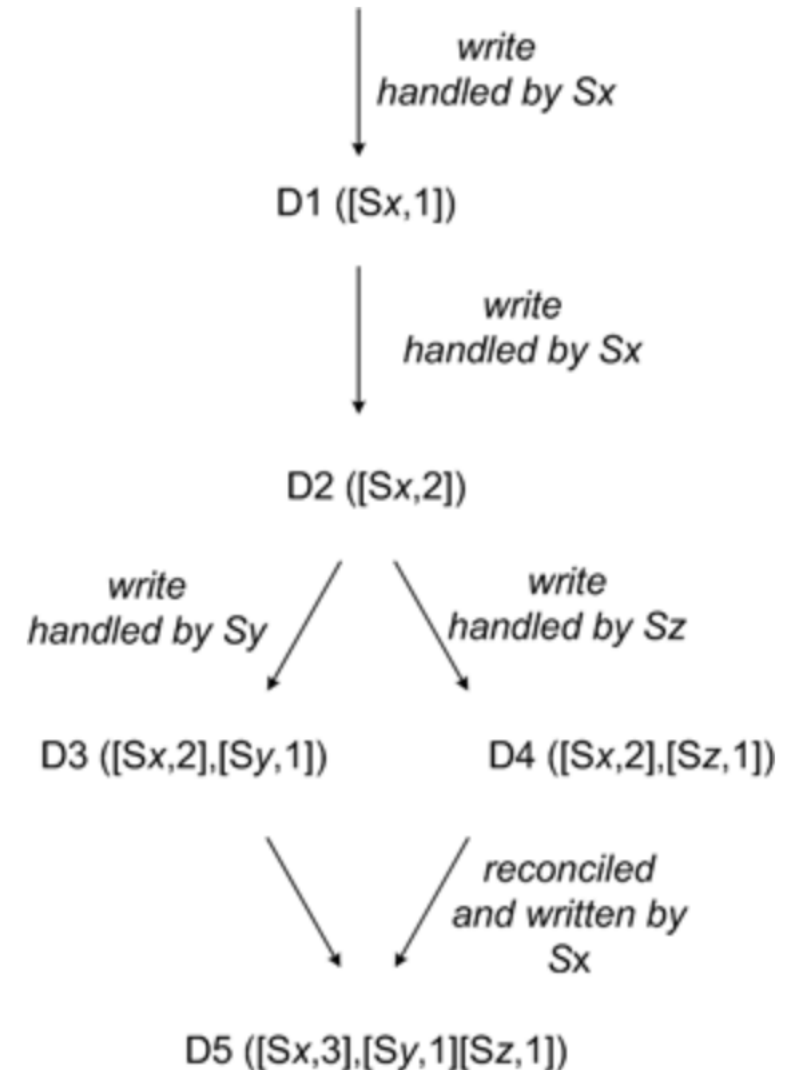
Source: Wiki
Source: MGH Co.

Example




Consistency using versions

- Vector clocks associated with every write
- Every read operation can understand the causal order and reconcile multiple versions
- Vector clocks can grow over time
- Dynamo maintains a timestamp when last time a node updated a data item and deletes old nodes beyond a limit




Handling faults - hinted handoff

- All reads and writes are performed by first N healthy nodes and not N nodes in exact sequence
 - ✓ This is called a “sloppy quorum”
 - ✓ Improves Availability when nodes fail
- Suppose A was supposed to receive a replica along with B and C but since A failed, it will go to D temporarily, along with a hint to D that A was the original choice
- Once A is up, D will send replica to A remove from itself
- Typically the key hash based replica nodes are kept across DCs to be highly durable
- Also typically ack for writes is kept > 1 so that at least 2 copies are written before ack to client




What about permanent failures of nodes ?

- Detect inconsistencies between replicas and solve it using minimum data transfers, e.g. if A is permanently down, how to make sure replicas are fixed on affected neighbors ?
- Use *Merkel trees* (hash trees) to track down subtrees where replicas have gone out of sync
 - ✓ Leaves are hashes of individual keys
 - ✓ Parents are hashes of children



Tweaking N, R, W for performance, durability, consistency, availability

- N - how many replicas total for a data item
 - R - how many replicas to return read success
 - W - how many replicas to return write success
-
- $W = 1$: High performance writes
 - $W = N$: Highly durable writes
 - $R = 1$: High performance reads
 - $R = N$: Highly consistent reads
-
- Typical config: $(N, R, W) = (3, 2, 2)$
 - Hard engineering problem to use commodity heterogeneous hardware and provide 99.9th percentile SLA on latency and accommodate slowest replicas for R and W targets



DynamoDB: Overview (1)

- NoSQL store that can store nested JSON documents in tables. Upto 32 levels of nesting.
- Partition key using consistent hashing
- Optional sorting key to sort within partition

People

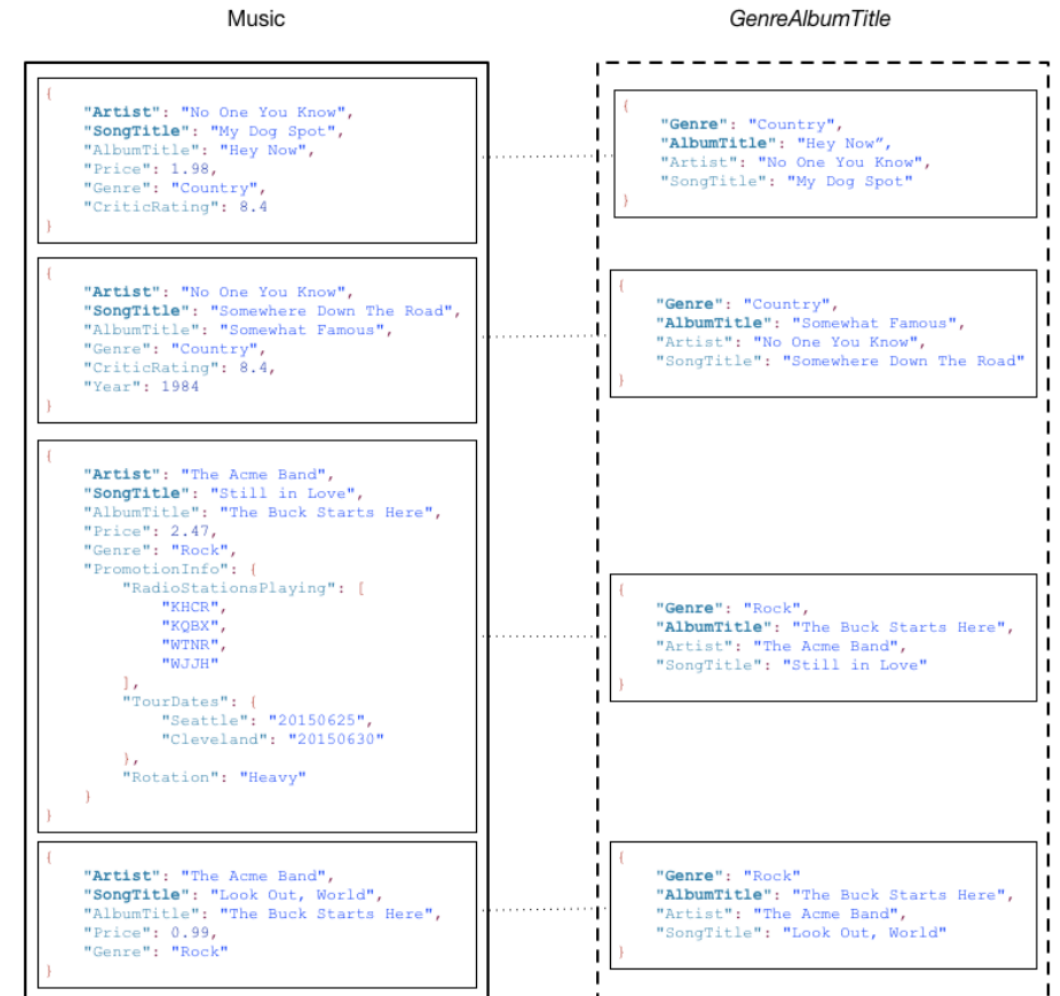
```
{  
  "PersonID": 101,  
  "LastName": "Smith",  
  "FirstName": "Fred",  
  "Phone": "555-4321"  
}
```

```
{  
  "PersonID": 102,  
  "LastName": "Jones",  
  "FirstName": "Mary",  
  "Address": {  
    "Street": "123 Main",  
    "City": "Anytown",  
    "State": "OH",  
    "ZIPCode": 12345  
  }  
}
```

```
{  
  "PersonID": 103,  
  "LastName": "Stephens",  
  "FirstName": "Howard",  
  "Address": {  
    "Street": "123 Main",  
    "City": "London",  
    "PostalCode": "ER3 5K8"  
  },  
  "FavoriteColor": "Blue"  
}
```

DynamoDB: Overview (2)

- Secondary indices
 - ✓ Global: Different from partition and sort keys
 - ✓ Local: Same partition but different from sort key



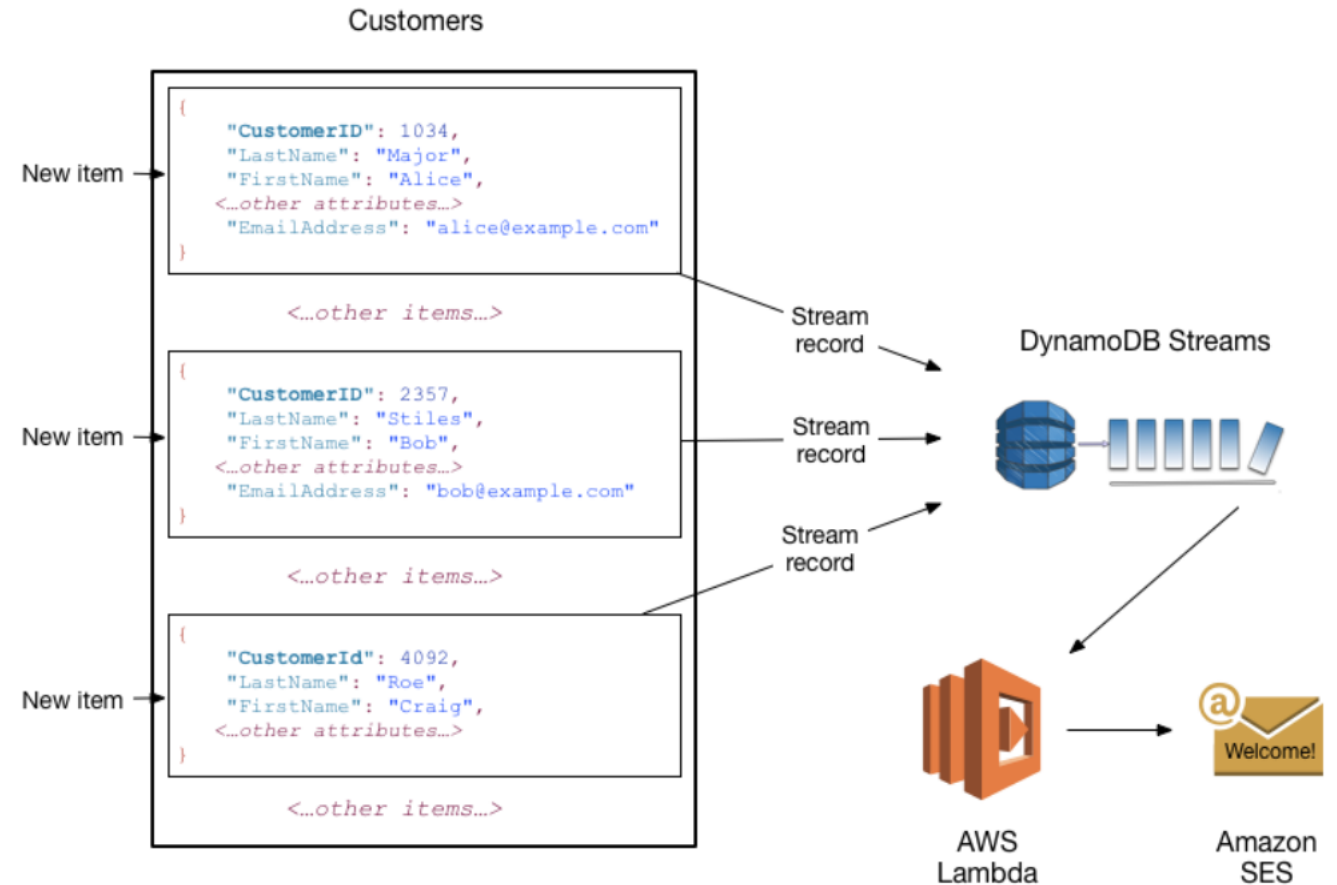
Global secondary index

DynamoDB: Overview (3)

- Streams

- ✓ Captures data modification event if enabled on a table
- ✓ On new item, update, delete the item (before+after for a change), timestamp, table name, meta-data create a Stream record

- Triggers can be created by adding a Lambda fn



e.g. send welcome email to new customers

DynamoDB: Consistency

- A table may be in multiple regions and replicated within region at AZs
- Eventual consistency reads
 - ✓ Faster reads
 - ✓ But stale data may be returned
 - ✓ Typically consistent within 1 sec
- Strong consistency (set ConsistentRead=True during query)
 - ✓ Returns latest write for reads
 - ✓ Reads may be unavailable on outages
 - ✓ Higher latency reads
 - ✓ Global secondary indices not supported for reads
 - ✓ Uses more throughput capacity for reads

DynamoDB: Read/Write Capacity modes

- Billing depends on how mode is set on a table
- On-demand
 - ✓ Single digit ms response on R/W
 - ✓ Needs no capacity planning
 - ✓ Good for unpredictable workloads that need fast response and pay-for-use
 - ✓ Auto-scaling happens within time limits. E.g. scales if workload doubles after 30min and throttles otherwise
- Provisioned (default)
 - ✓ Reserve capacity for read / write throughput needed
 - ✓ Cost predictability
 - ✓ Can still auto-scale between levels
- Capacity units
 - ✓ Read request unit = 1 x Strongly consistent read or 2 x Eventually consistent reads of 4KB items
 - ✓ Write request unit = 1 x 1 KB write or 0.5 x transactional 1KB write

DynamoDB: CRUD

```
{
  TableName : "Music",
  KeySchema: [
    {
      AttributeName: "Artist",
      KeyType: "HASH", //Partition key
    },
    {
      AttributeName: "SongTitle",
      KeyType: "RANGE" //Sort key
    }
  ],
  AttributeDefinitions: [
    {
      AttributeName: "Artist",
      AttributeType: "S"
    },
    {
      AttributeName: "SongTitle",
      AttributeType: "S"
    }
  ],
  ProvisionedThroughput: { // Only specified if using provisioned mode
    ReadCapacityUnits: 1,
    WriteCapacityUnits: 1
  }
}
```

```
INSERT INTO Music
  (Artist, SongTitle, AlbumTitle,
   Year, Price, Genre,
   Tags)
VALUES(
  'No One You Know', 'Call Me Today', 'Somewhat Famous',
  2015, 2.14, 'Country',
  '{"Composers": ["Smith", "Jones", "Davis"],"LengthInSeconds": 214}'
);
```

```
/* Return all of the songs by an artist, with a particular word in the title...
...but only if the price is less than 1.00 */
```

```
SELECT * FROM Music
WHERE Artist='No One You Know' AND SongTitle LIKE '%Today%'
AND Price < 1.00;
```

```
// Return all of the songs by an artist, matching first part of title
```

```
{
  TableName: "Music",
  KeyConditionExpression: "Artist = :a and begins_with(SongTitle, :t)",
  ExpressionAttributeValues: {
    ":a": "No One You Know",
    ":t": "Call"
  }
}
```

- Options
 - ✓ PartiQL
 - ✓ APIs



**Next Session:
Spark introduction**