



**BITS
Pilani**

Pilani Campus

Deep Learning

Dr. Sugata Ghosal

sugata.ghosal@pilani.bits-pilani.ac.in

Agenda



- Introduction
- Course Objectives and Logistics
- Introduction to Perceptron and MLP
 - Approximation Capabilities
- Characteristics of Deep Learning
- Reading: Chapter 1 of Textbook



Neural Networks are taking over!

- Neural networks have become one of the major thrust areas recently in various pattern recognition, prediction, and analysis problems
- In many problems they have established the state of the art
 - Often exceeding previous benchmarks by large margins

Breakthroughs with neural networks



www.technewsworld.com/story/84013.html

TECHNEWSWORLD EMERGING TECH

Computing Internet IT Mobile Tech Reviews Security Technology Tech Blog Reader Services

Microsoft AI Beats Humans at Speech Recognition

By Richard Adhikari
Oct 20, 2016 11:40 AM PT

Print
Email




Image: Adobe Stock

Microsoft's Artificial Intelligence and Research Unit earlier this week reported that its speech recognition technology had surpassed the performance of human transcriptionists.

How do you feel about Black Friday and Cyber Monday?

- ☐ They're great -- I get a lot of bargains!
- ☐ The deals are too spread out -- I'd prefer just one day.
- ☐ They're a fun way to kick off the holiday season.
- ☐ I don't like the commercialization of Thanksgiving Day.
- ☐ They're crucial for the retail industry and the economy.
- ☐ The deals typically aren't that good.

Vote to See Results

E-Commerce Times

Black Friday Shoppers Hungry for New Experiences, New Tech

Pay TV's Newest Innovation: Giving Users Control

Apple Celebrates Itself in \$300 Coffee Table Tome

AWS Enjoys Top Perch in IaaS, PaaS Markets

US Comptroller Gears Up for Blockchain and

Breakthrough with neural networks

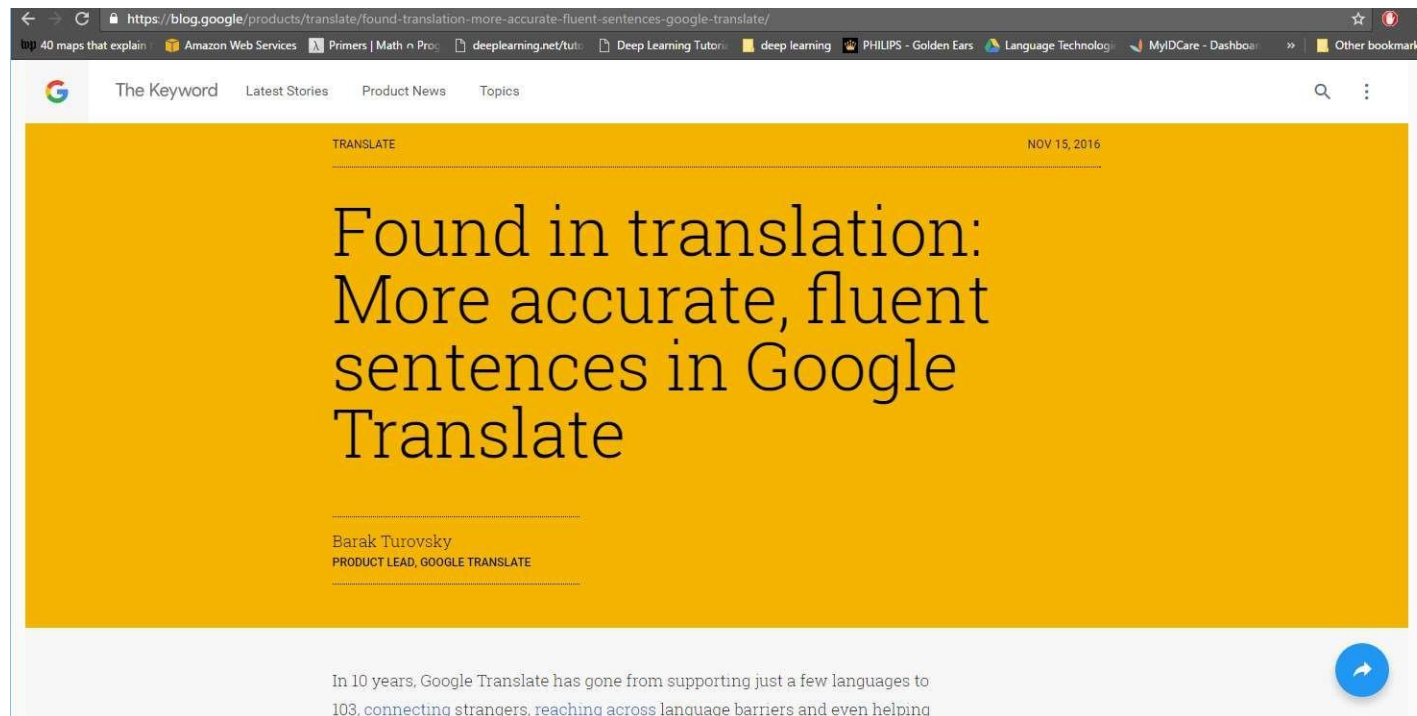


Image segmentation and recognition





Image recognition



<https://www.sighthound.com/technology/>

Breakthroughs with neural networks


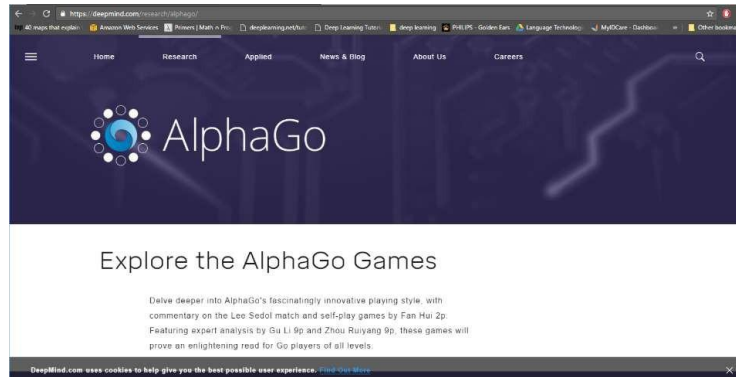
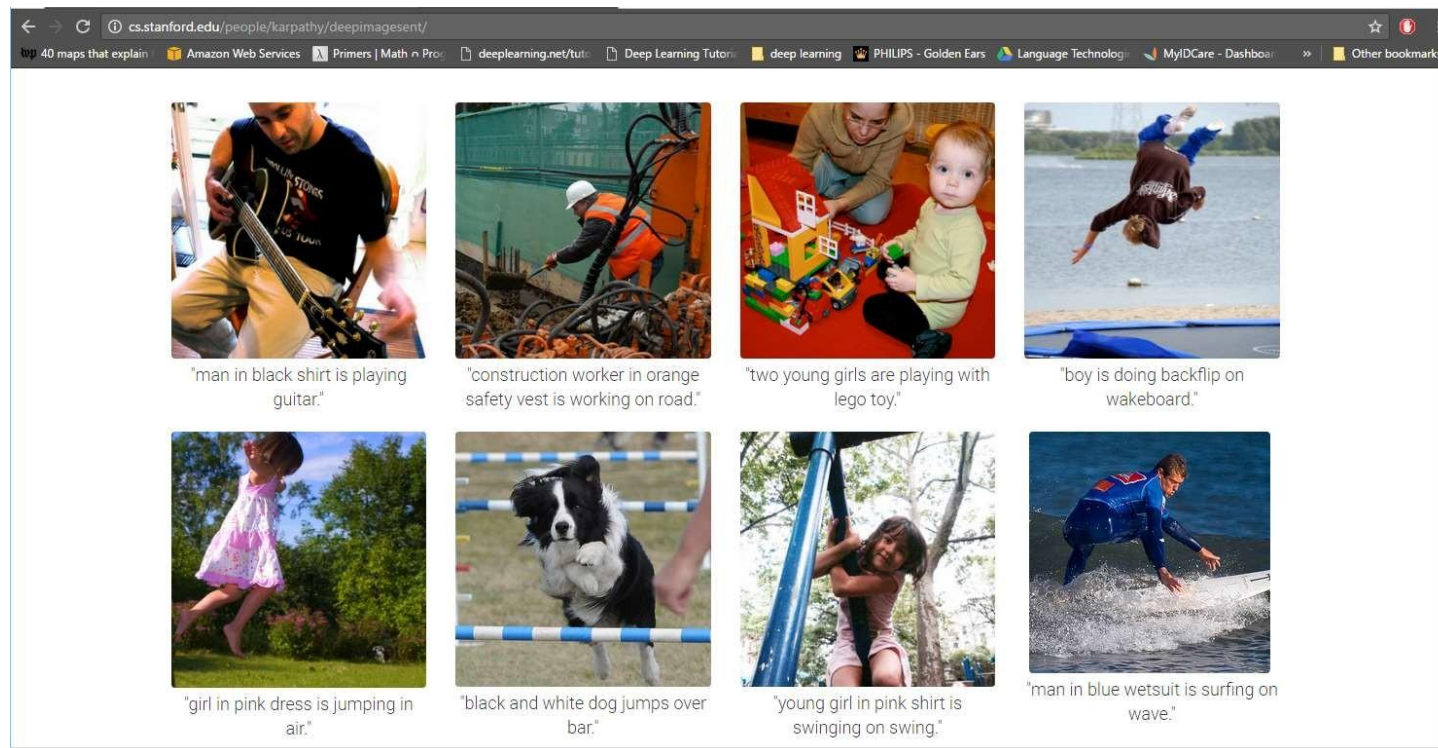


Figure 1: Training *AlphaZero* for 700,000 steps. Elo ratings were computed from evaluation games between different players when given one second per move. **a** Performance of *AlphaZero* in chess, compared to 2016 TCEC world-champion program *Stockfish*. **b** Performance of *AlphaZero* in shogi, compared to 2017 CSA world-champion program *Elmo*. **c** Performance of *AlphaZero* in Go, compared to *AlphaGo Lee* and *AlphaGo Zero* (20 block / 3 day) (29).

Success with neural networks



- Captions generated entirely by a neural network

Successes with neural networks



- And a variety of other problems:
 - From art to astronomy to healthcare..
 - and even predicting stock markets!
-



Objectives of this course

- Understanding neural networks
 - Comprehending the models that do the previously mentioned tasks
 - And maybe build them
 - Design, build and train networks for various tasks
 - *You will not become an expert in one course*
-

Course objectives: Broad level



Deep Dive into Artificial Neural Networks

- Concepts
 - Types of neural networks and underlying ideas
 - Learning in neural networks
 - Training, concepts, practical issues
 - Architectures and applications
 - Practical
 - Familiarity with training and parameter tuning
 - Implement various neural network architectures
 - Overall: Set you up for further work in your area
-

Course learning objectives:

Topics



- Basic network formalisms (for classification and prediction):
 - Multi-Layer Perceptron (MLP)
 - Convolutional networks (CNN)
 - Recurrent networks (RNN)
- Some advanced formalisms (for creation)
 - Generative models: VAEs
 - Adversarial models: GANs
- Applications we will touch upon:
 - Computer vision: recognizing images
 - Text processing: modelling and generating language
 -



Reading

- List of books on Canvas Course Page
 - Primary: <https://www.deeplearningbook.org/>
 - “Deep Learning”, Goodfellow, Bengio, Courville
 - Reference:
<https://www.manning.com/books/deep-learning-with-python>
 - “Deep Learning with Python”, Francois Chollet.
 - Additional reading material will be posted on Canvas, if needed
-



Logistics

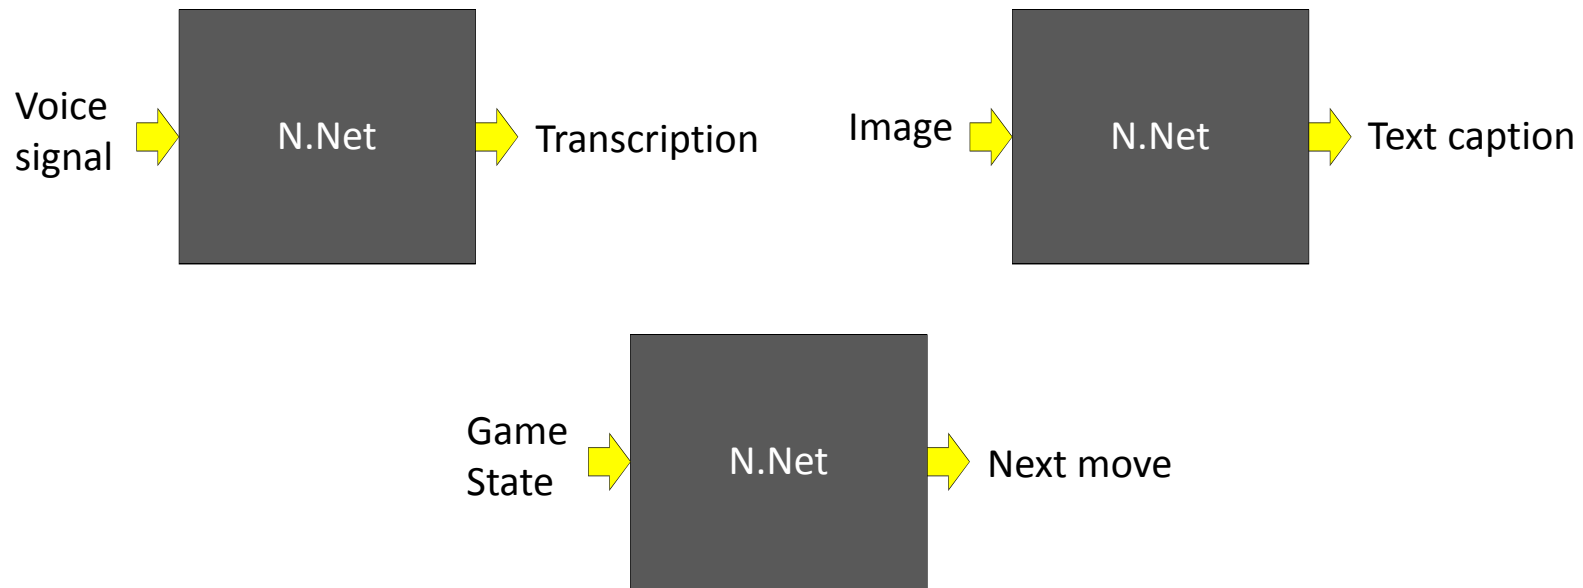
- Most relevant info on Canvas
 - Handout
 - Schedule of Webinars, Quiz, Assignments,
 - Lecture Slides
 - Lab Sheets
 - One Quiz, Two Assignments
 - Quiz, one assignment before midsem
 - One assignment after midsem
 - *submissions beyond deadline will be deducted some marks / day (unless medical emergencies)*
 - Programming using Python, Keras / Tensorflow
-



Questions?

- Please post on Discussions Forum
 - TAs and instructors will answer
 - Collaborate with your fellow students

So what are neural networks??



- What are these boxes?

So what are neural networks??



- It begins with this..
-

Early Models of Human Cognition



- Associationism
 - Humans learn through association
 - 400BC-1900AD: Plato, David Hume, Ivan Pavlov..
-

Observation: *The Brain*




- Mid 1800s: The brain is a mass of interconnected neurons
-

Brain: Interconnected Neurons



- Many neurons connect *in* to each neuron
 - Each neuron connects *out* to many neurons
-


Connectionist Machines



- Network of processing elements
- **All world knowledge is stored in the *connections* between the elements**
 - *But what are these individual elements?*


Modelling the brain

- What are the units?
- A neuron:



- Signals come in through the dendrites into the Soma
- A signal goes out via the axon to other neurons
 - Only one axon per neuron
- Factoid that may only interest me: Neurons do not undergo cell division
 - Neurogenesis occurs from neuronal stem cells, and is minimal after birth

Rosenblatt's perceptron




- Original perceptron model
 - Groups of sensors (S) on retina combine onto cells in association area A_I
 - Groups of A_I cells combine into Association cells A_{II}
 - Signals from A_{II} cells combine into response cells R
 - All connections may be excitatory or inhibitory

Connectionist Machines




- Neural networks are *connectionist* machines
 - As opposed to Von Neumann Machines



- The machine has many non-linear processing units
 - The program is the connections between these units
 - Connections may also define memory

Simplified mathematical model of Perceptron



- Number of inputs combine linearly
 - Threshold logic: Fire if combined input exceeds or equal to threshold

$$Y = \begin{cases} 1 & \text{if } \sum_i w_i x_i - T \geq 0 \\ 0 & \text{else} \end{cases}$$

Also provided a learning algorithm



$$\mathbf{w} = \mathbf{w} + \eta(d(\mathbf{x}) - y(\mathbf{x}))\mathbf{x}$$

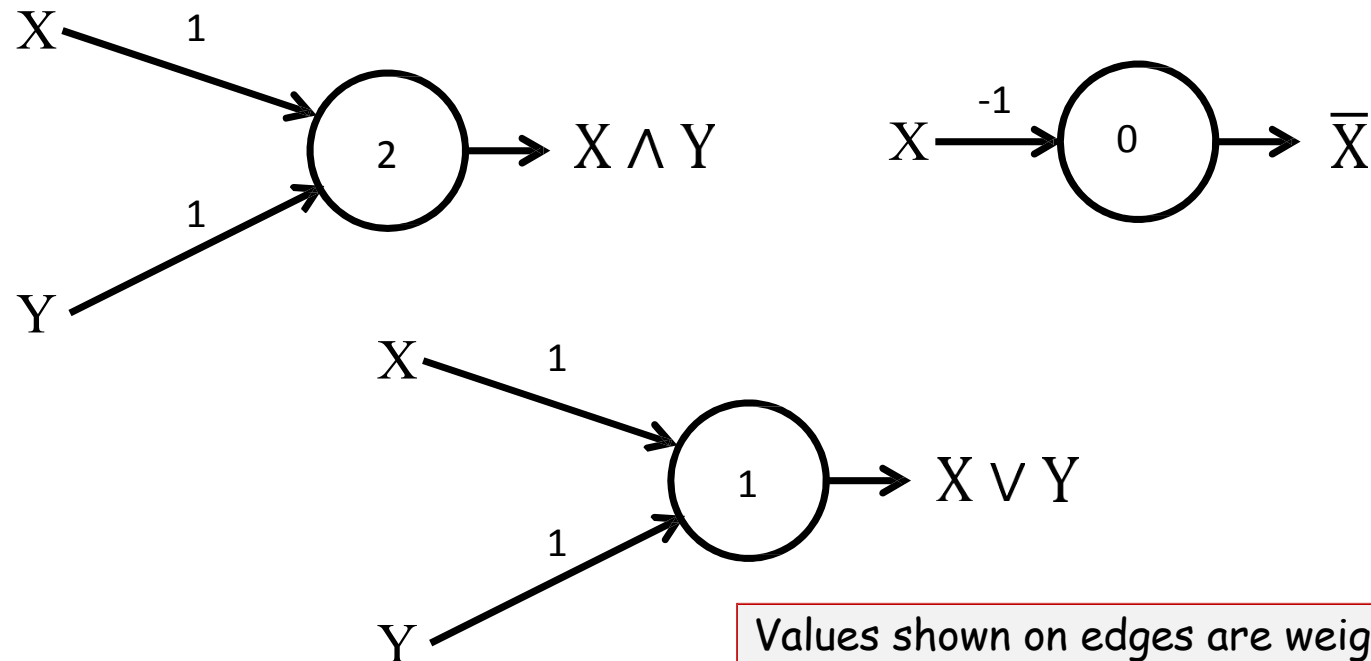
Sequential Learning:

$d(x)$ is the desired output in response to input x

$y(x)$ is the actual output in response to x

- Boolean tasks
 - Update the weights whenever the perceptron output is wrong
 - Proved convergence for linearly separable classes
-

Perceptron




Values shown on edges are weights,
numbers in the circles are thresholds

- Easily shown to mimic any Boolean gate
- But...


Perceptron

No solution for XOR!
Not universal!




- Minsky and Papert, 1968
-

A single neuron is not enough



- Individual elements are weak computational elements
 - Marvin Minsky and Seymour Papert, 1969,
Perceptrons: An Introduction to Computational Geometry
- *Networked* elements are required

Multi-layer Perceptron!




- **XOR**

- The first layer is a “hidden” layer
- Also originally suggested by Minsky and Papert 1968


A more generic model

$$((A \& \bar{X} \& Z) | (A \& \bar{Y})) \& ((X \& Y) | \overline{(X \& Z)})$$




- A “multi-layer” perceptron
- Can compose arbitrarily complicated Boolean functions!
 - In cognitive terms: Can compute arbitrary Boolean functions over sensory input
 - More on this in the next class

But our brain is not Boolean




- We have real inputs
 - We make non-Boolean inferences/predictions
-

The perceptron with *real* inputs



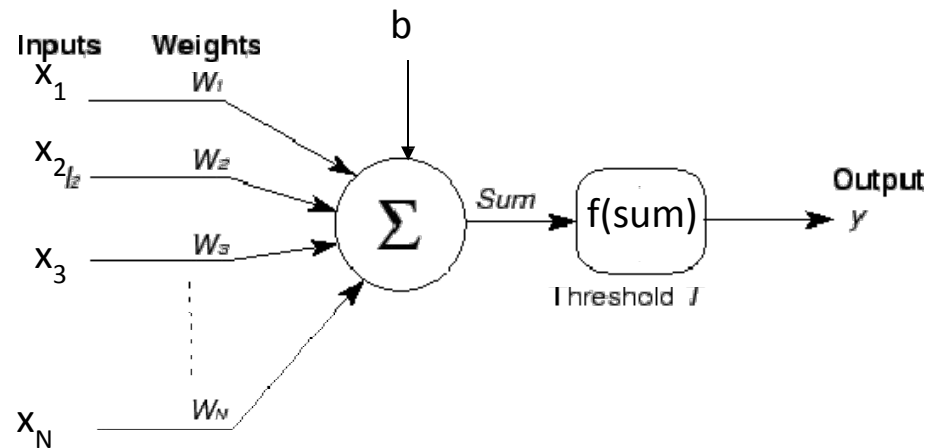
- $x_1 \dots x_N$ are real valued
- $w_1 \dots w_N$ are real valued
- Unit “fires” if weighted input exceeds a threshold

The perceptron with *real* inputs **and a real output**




- $x_1 \dots x_N$ are real valued
- $w_1 \dots w_N$ are real valued
- The output y can also be real valued
 - Sometimes viewed as the “probability” of firing

The “real” valued perceptron



- Any real-valued “activation” function may operate on the weighted- sum input
 - We will see several later
 - Output will be real valued
- The perceptron maps real-valued inputs to real-valued outputs
- *Is useful to continue assuming Boolean outputs though, for interpretation*



A Perceptron on Reals




$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i \geq T \\ 0 & \text{else} \end{cases}$$

- A perceptron operates on *real-valued* vectors

– This is a **linear classifier**




Boolean functions with a real perceptron



- Boolean perceptrons are also linear classifiers
 - Purple regions have output 1 in the figures
 - What are these functions
 - Why can we not compose an XOR?


Composing complicated “decision” boundaries



Can now be composed into
“networks” to compute arbitrary
classification “boundaries”


- Build a network of units with a single output that fires if the input is in the coloured area

Booleans over the reals




- The network must fire if the input is in the coloured area

Booleans over the reals





- The network must fire if the input is in the coloured area

Booleans over the reals





- The network must fire if the input is in the coloured area

Booleans over the reals




- The network must fire if the input is in the coloured area

Booleans over the reals




- The network must fire if the input is in the coloured area

Booleans over the reals




- The network must fire if the input is in the coloured area

More complex decision boundaries




- Network to fire if the input is in the yellow area
 - “OR” two polygons
 - A third layer is required

Complex decision boundaries



- Can compose very complex decision boundaries
 - How complex exactly? More on this in the next class
-

Complex decision boundaries




- Classification problems: finding decision boundaries in high-dimensional space
 - Can be performed by an MLP
- MLPs can *classify* real-valued inputs



Story so far


- **MLPs are connectionist computational models**
 - Individual perceptrons are computational equivalent of neurons
 - The MLP is a layered composition of many perceptrons
 - **MLPs can model Boolean functions**
 - Individual perceptrons can act as Boolean gates
 - Networks of perceptrons are Boolean functions
 - **MLPs are Boolean *machines***
 - They represent Boolean functions over linear boundaries
 - They can represent arbitrary decision boundaries
 - They can be used to *classify* data
-

But what about continuous valued *output*?





- Inputs may be real valued
- Can outputs be continuous-valued too?

MLP as a continuous-valued regression



- A simple 3-unit MLP with a “summing” output unit can generate a “square pulse” over an input
 - Output is 1 only if the input lies between T_1 and T_2
 - T_1 and T_2 can be arbitrarily specified

MLP as a continuous-valued regression



- A simple 3-unit MLP can generate a “square pulse” over an input
- **An MLP with many units can model an arbitrary function over an input**
 - To arbitrary precision
 - Simply make the individual pulses narrower
- This generalizes to functions of any number of inputs

Story so far




- **Multi-layer perceptrons are connectionist computational models**
- **MLPs are *classification engines***
 - They can identify classes in the data
 - Individual perceptrons are feature detectors
 - The network will fire if the combination of the detected basic features matches an “acceptable” pattern for a desired class of signal
- **MLP can also model continuous valued functions**

So what does the perceptron really model?



- Is there a “semantic” interpretation?
 - Cognitive version: Is there an interpretation beyond the simple characterization as Boolean functions over sensory inputs?
-

Lets look at the weights




$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i \geq T \\ 0 & \text{else} \end{cases}$$

$$y = \begin{cases} 1 & \text{if } \mathbf{x}^T \mathbf{w} \geq T \\ 0 & \text{else} \end{cases}$$

- What do the *weights* tell us?
 - The neuron fires if the inner product between the weights and the inputs exceeds a threshold


The weight as a “template”



$$X^T W > T$$


$$\cos \theta > \frac{T}{|X|}$$

$$\theta < \cos^{-1} \left(\frac{T}{|X|} \right)$$




- The perceptron fires if the input is within a specified angle of the weight
- Neuron fires if the input vector is close enough to the weight vector.
 - If the input pattern matches the weight pattern closely enough


The weight as a template



1 If $\sum w_i x_i \geq 0$,
0 otherwise



Correlation = 0.57

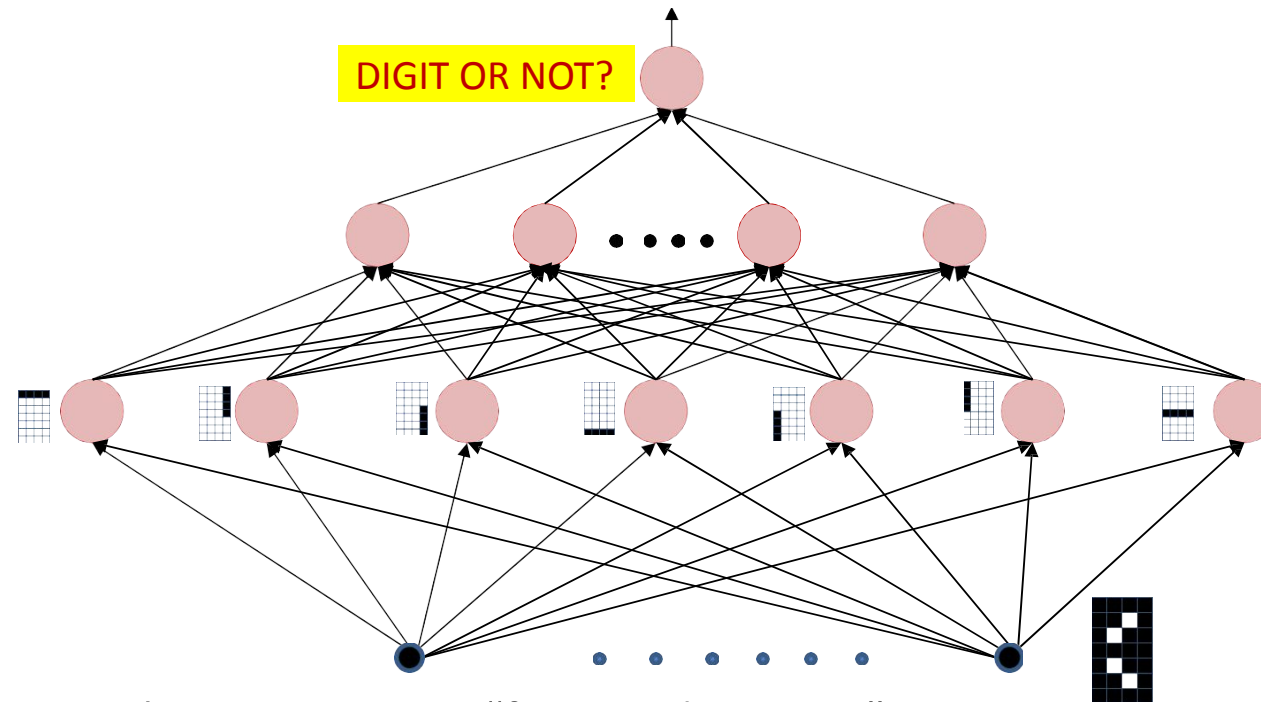


Correlation = 0.82




- If the *correlation* between the weight pattern and the inputs exceeds a threshold, fire
- The perceptron is a *correlation filter*!

The MLP as a Boolean function over feature detectors



- The input layer comprises “feature detectors”
 - Detect if certain patterns have occurred in the input
- The network is a Boolean function over the feature detectors
- I.e. it is important for the *first* layer to capture relevant patterns

The MLP as a cascade of feature detectors



- The network is a cascade of feature detectors
 - Higher level neurons compose complex templates from features represented by lower-level neurons

Terminology

